



# Telco Customer Churn Analysis



Author: Muhammad Tahir



Date: 30 March 2025



## What I did here:

In this notebook, I explored the Telco Customer Churn dataset, broke down key trends and factors causing churn.



## Importing Libraries and Dataset

### 1.1 Load necessary libraries

```
In [4]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

### 1.2 Import dataset

```
In [6]: df = pd.read_csv('Telco_Customer_Churn.csv')
```

### 1.3 Quick dataset overview

```
In [8]: df.sample(5)
```

```
Out[8]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	Phone
<b>3785</b>	8337-UPOAQ	Male	1	Yes	No	10	
<b>4082</b>	0112-QWPNC	Male	0	Yes	No	49	
<b>2679</b>	7341-LXCAF	Male	0	Yes	No	4	
<b>6760</b>	5295-PCJOO	Male	0	No	Yes	4	
<b>1163</b>	0135-NMXAP	Female	0	No	No	12	

5 rows × 21 columns

## 2 Data Understanding & Basic Checks

### 2.1 Dataset shape

```
In [11]: df.shape
```

```
Out[11]: (7043, 21)
```

### 2.2 Check columns names and dtypes

```
In [13]: df.info(verbose=True)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

## 2.3 Check for missing values

```
In [15]: df.isnull().sum()
```

```

Out[15]: customerID            0
gender                0
SeniorCitizen         0
Partner               0
Dependents            0
tenure                0
PhoneService          0
MultipleLines         0
InternetService       0
OnlineSecurity        0
OnlineBackup          0
DeviceProtection      0
TechSupport           0
StreamingTV           0
StreamingMovies       0
Contract              0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges          0
Churn                 0
dtype: int64

```

## 2.4 Quick statistical summary

```
In [17]: df.describe()
```

```
Out[17]:
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

- The **SeniorCitizen** column represents a categorical variable (0 = No, 1 = Yes), meaning percentile-based distribution (25%-50%-75%) isn't meaningful for analysis. Only **16.2%** of the customers fall into this category.
- Regarding **tenure**, 75% of customers have been subscribed for **less than 55 months**, with an average tenure of **32 months**. However, some customers have remained for as long as **72 months** (6 years), showing significant variability.
- For monthly charges, the average bill stands at 64.76, *but pricing varies widely. While half of the customers pay 70.35 or less, the top 25% are billed over 89.85 per month. The highest monthly charge reaches 118.75*, which is significantly above the average.

## 2.5 Check Duplicates

```
In [20]: df.duplicated().sum()
```

```
Out[20]: 0
```

## 3 Data Cleaning & Preprocessing

```
In [22]: # TotalCharges is stored as an object; it needs conversion
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
```

```
In [23]: df.isnull().sum()
```

```
Out[23]: customerID      0
gender      0
SeniorCitizen  0
Partner     0
Dependents  0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport  0
StreamingTV  0
StreamingMovies  0
Contract     0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  11
Churn        0
dtype: int64
```

```
In [24]: df.dropna(how='any', inplace=True)
```

```
In [25]: # convert 'SeniorCitizen' to categorical
df['SeniorCitizen'] = df['SeniorCitizen'].map({1:'Yes', 0:'No'})
```

```
In [26]: df.info(verbose=True)
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 7032 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7032 non-null   object
1   gender                 7032 non-null   object
2   SeniorCitizen          7032 non-null   object
3   Partner                7032 non-null   object
4   Dependents             7032 non-null   object
5   tenure                 7032 non-null   int64
6   PhoneService           7032 non-null   object
7   MultipleLines          7032 non-null   object
8   InternetService        7032 non-null   object
9   OnlineSecurity         7032 non-null   object
10  OnlineBackup           7032 non-null   object
11  DeviceProtection       7032 non-null   object
12  TechSupport            7032 non-null   object
13  StreamingTV            7032 non-null   object
14  StreamingMovies        7032 non-null   object
15  Contract               7032 non-null   object
16  PaperlessBilling       7032 non-null   object
17  PaymentMethod          7032 non-null   object
18  MonthlyCharges         7032 non-null   float64
19  TotalCharges           7032 non-null   float64
20  Churn                  7032 non-null   object
dtypes: float64(2), int64(1), object(18)
memory usage: 1.2+ MB

```

**Categorize customers into groups based on their tenure. For example, assign a tenure group of 1-12 for customers with a tenure of less than 12 months, 13-24 for those with a tenure between 1 and 2 years, and continue grouping in a similar manner.**

```
In [28]: df['tenure'].max()
```

```
Out[28]: 72
```

```
In [29]: # group tenure into bins of 12 months and assign labels to these groups
labels = [f"{i} - {i+12}" for i in range(1, 72, 12)]

df['tenure_group'] = pd.cut(df['tenure'], range(1, 80, 12), right=False, labels=labels)
```

```
In [30]: df['tenure_group'].value_counts()
```

```
Out[30]: tenure_group
1 - 12      2175
61 - 72     1407
13 - 24     1024
25 - 36      832
49 - 60      832
37 - 48      762
Name: count, dtype: int64
```

```
In [31]: # drop unnecessary columns
df.drop(columns=['customerID'], axis=1, inplace=True)
df.sample(5)
```

```
Out[31]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Mu
<b>6974</b>	Female	No	Yes	No	51	Yes	
<b>2331</b>	Female	No	Yes	No	25	Yes	
<b>4162</b>	Female	No	Yes	Yes	72	No	
<b>2192</b>	Male	Yes	Yes	No	56	Yes	
<b>6490</b>	Male	No	No	No	26	Yes	

5 rows × 21 columns

## Data Exploration

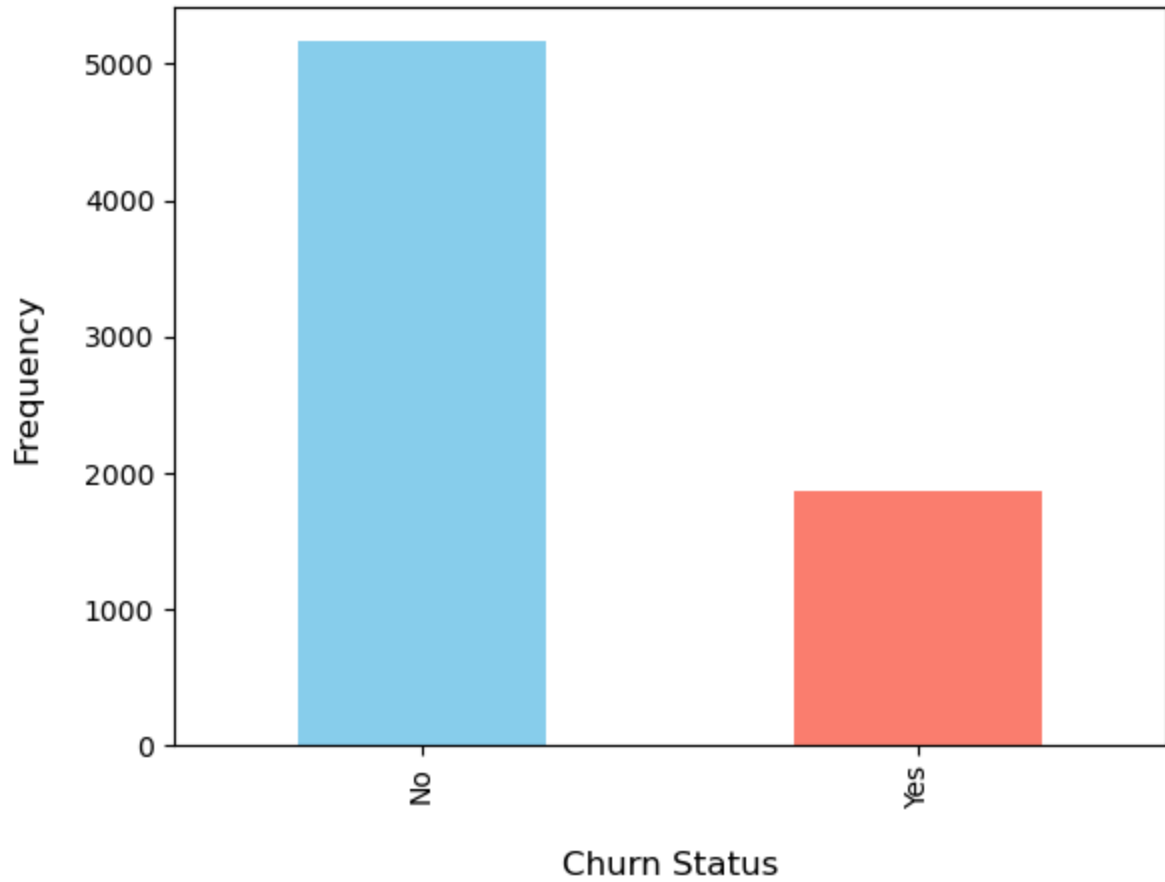
### Univariate Analysis

```
In [34]: df['Churn'].value_counts()
```

```
Out[34]: Churn
No      5163
Yes     1869
Name: count, dtype: int64
```

```
In [35]: # visulaizing the distribution of churn variable
df['Churn'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.ylabel('Frequency', fontsize=12, labelpad=15)
plt.xlabel('Churn Status', fontsize=12, labelpad=15)
plt.title('Distribution of Churn Categories', fontsize=14, fontweight='bold')
plt.show()
```

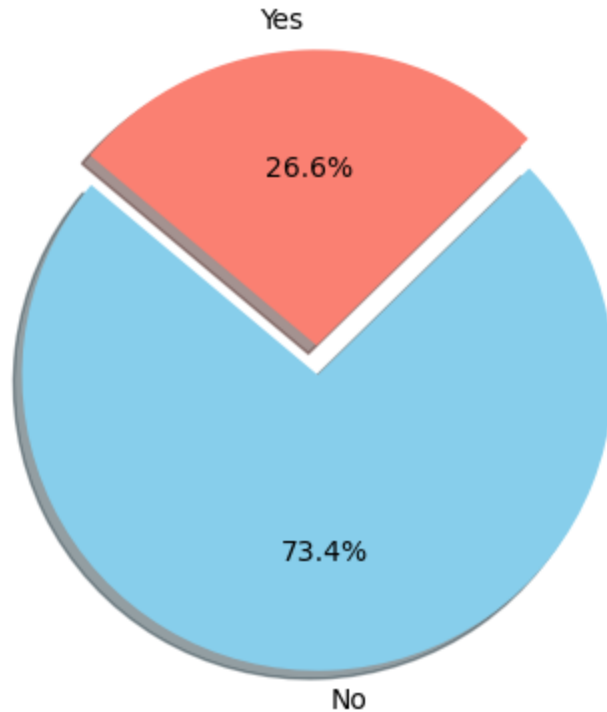
## Distribution of Churn Categories



```
In [36]: df['Churn'].value_counts().plot(kind='pie', autopct='%1.1f%%', startangle=14,
plt.ylabel("")
plt.title("Churn Rate Distribution", fontsize=14, fontweight='bold', pad=15)
plt.show()
```

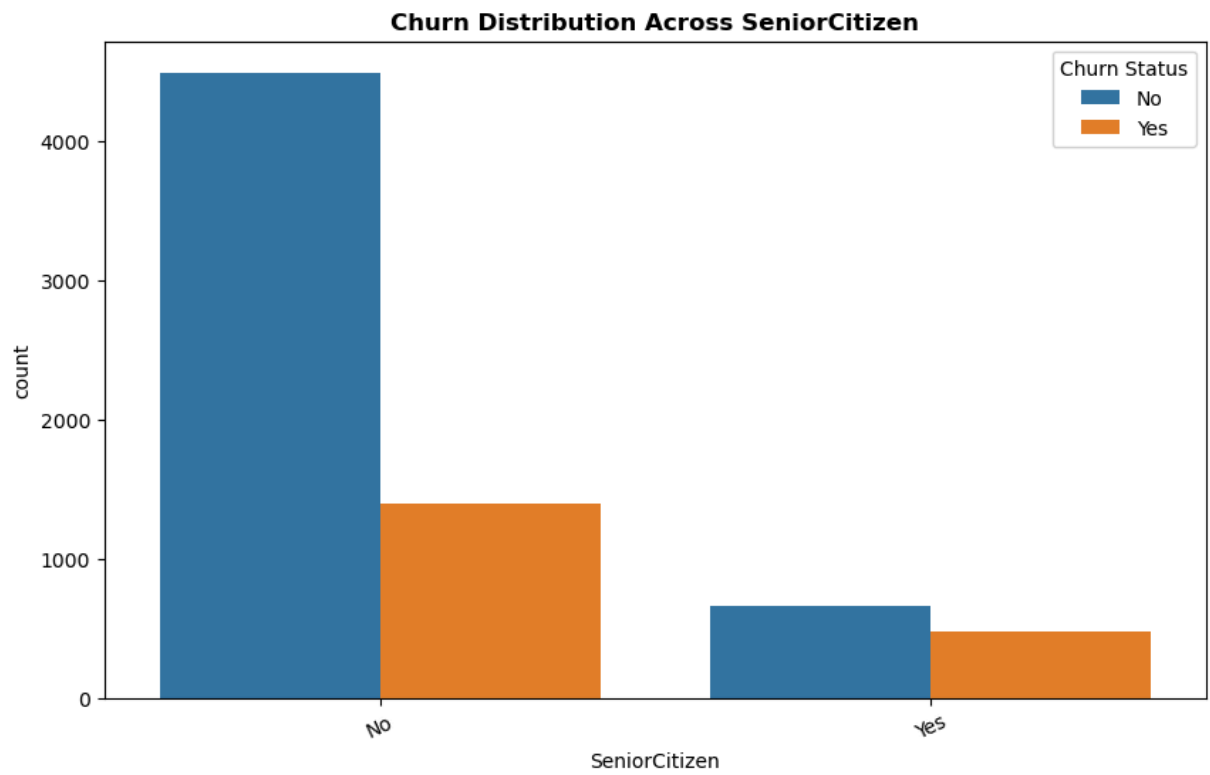
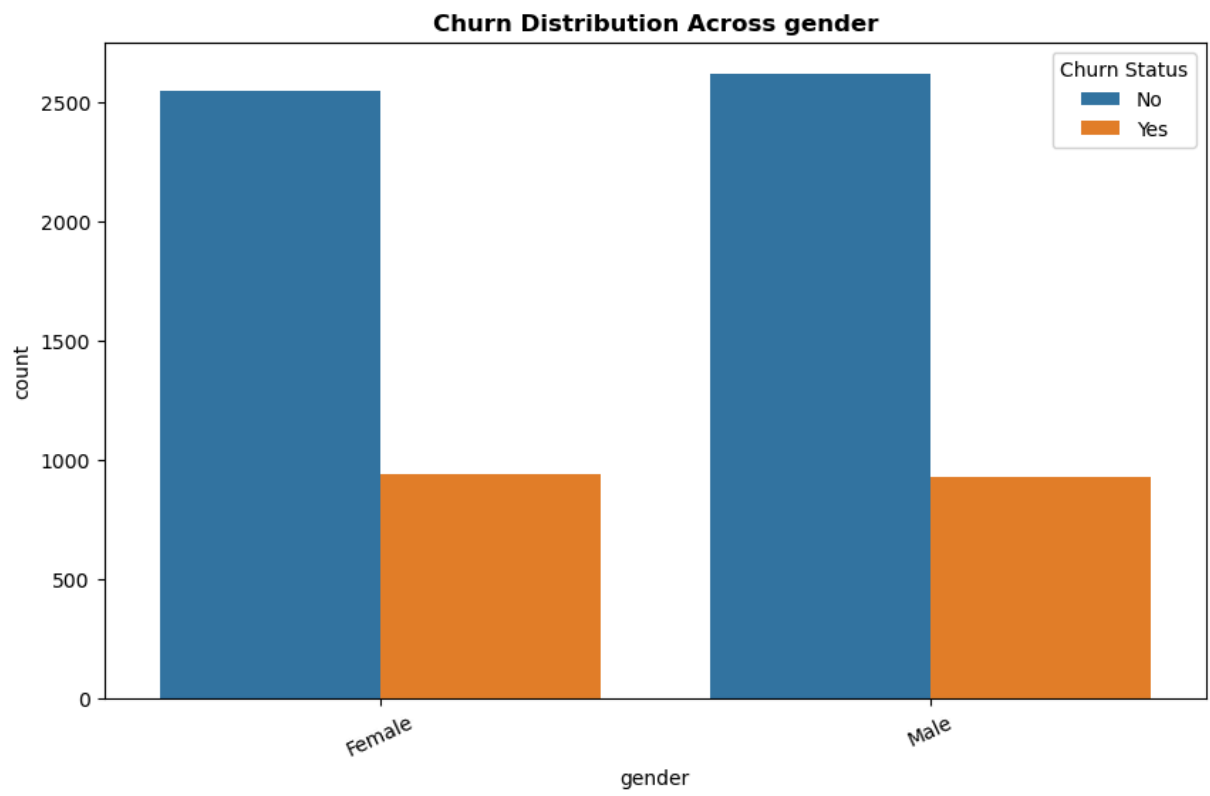


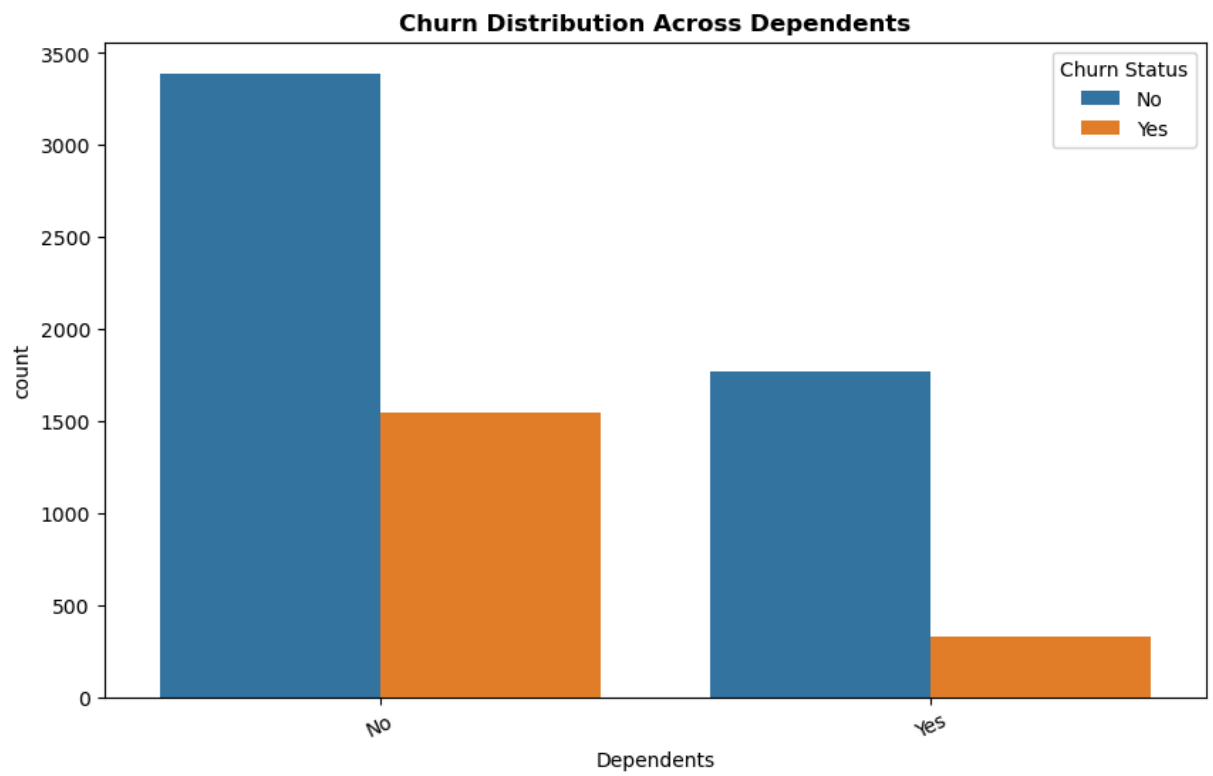
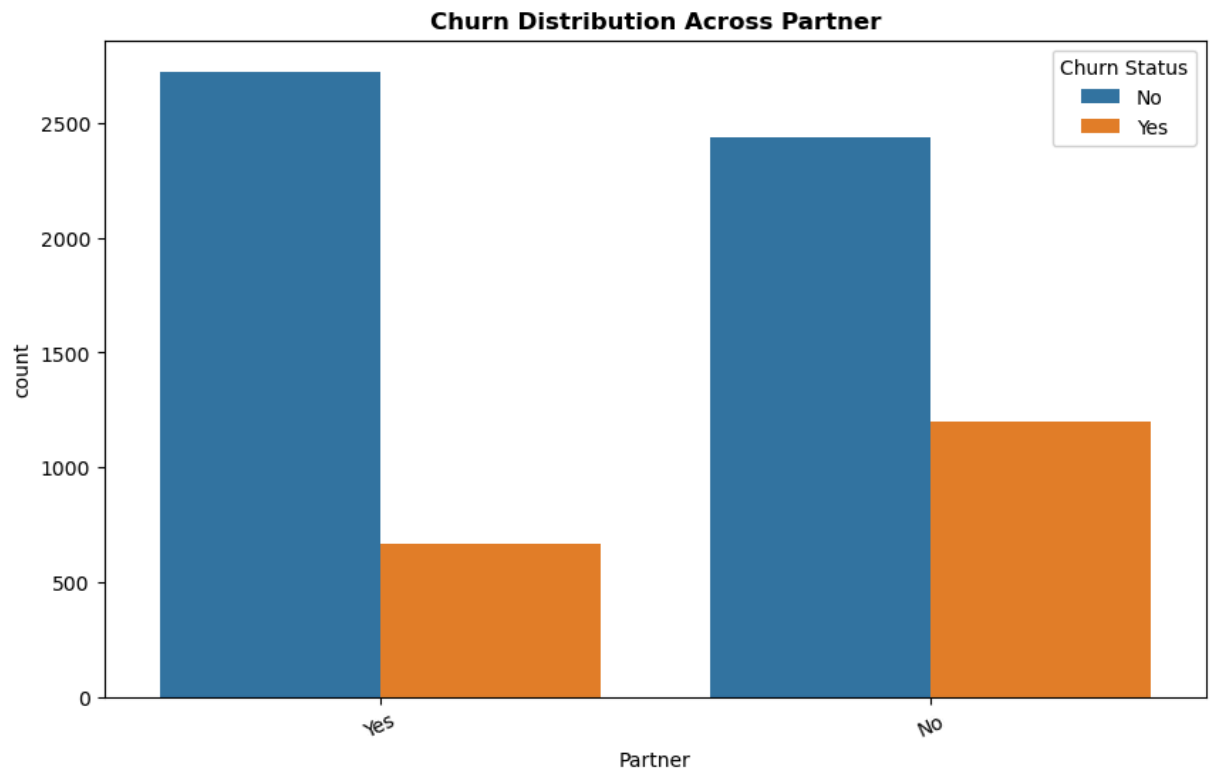
## Churn Rate Distribution

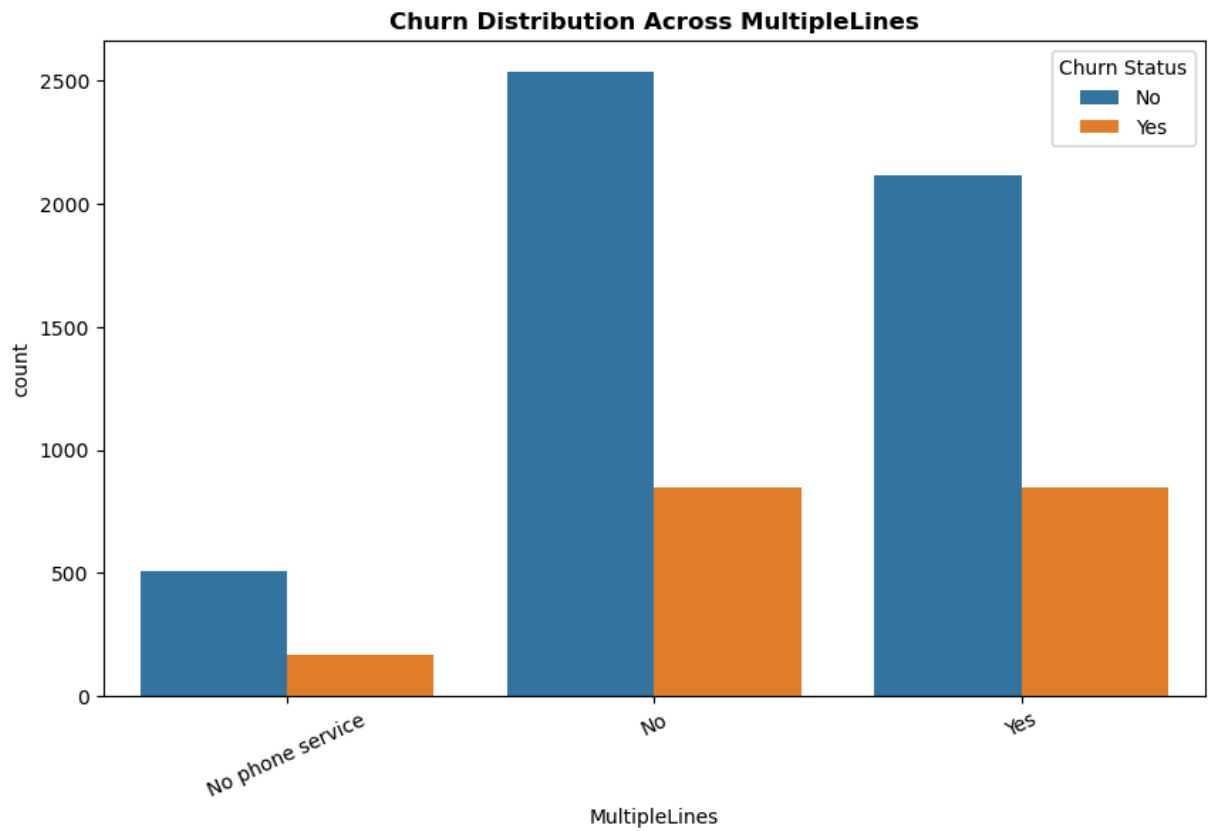
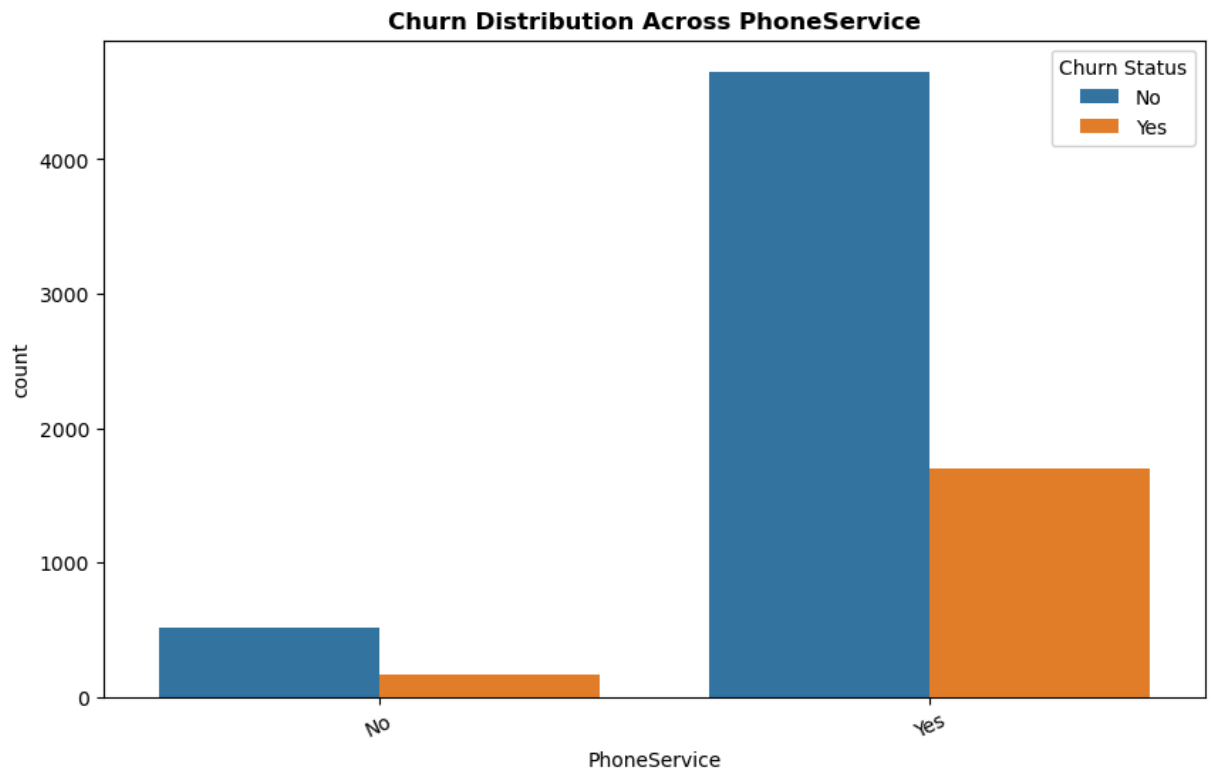


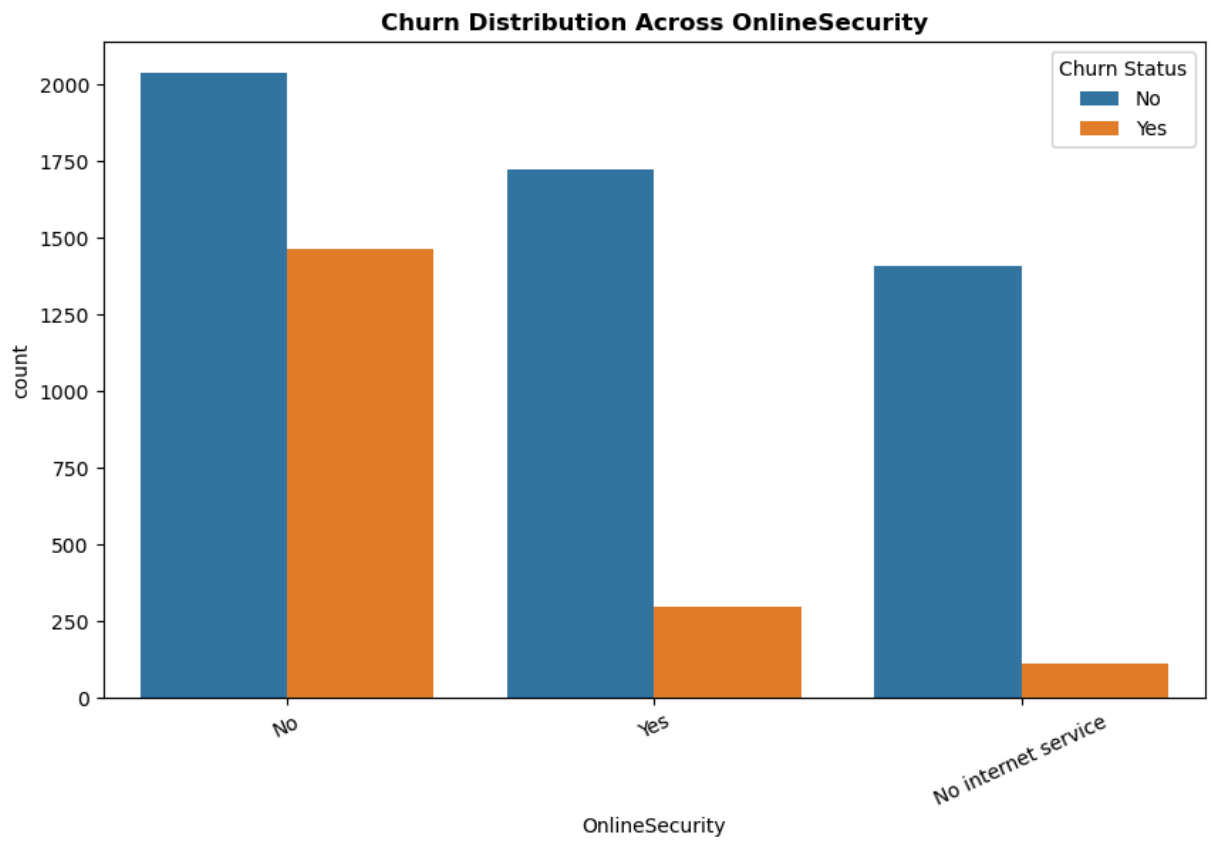
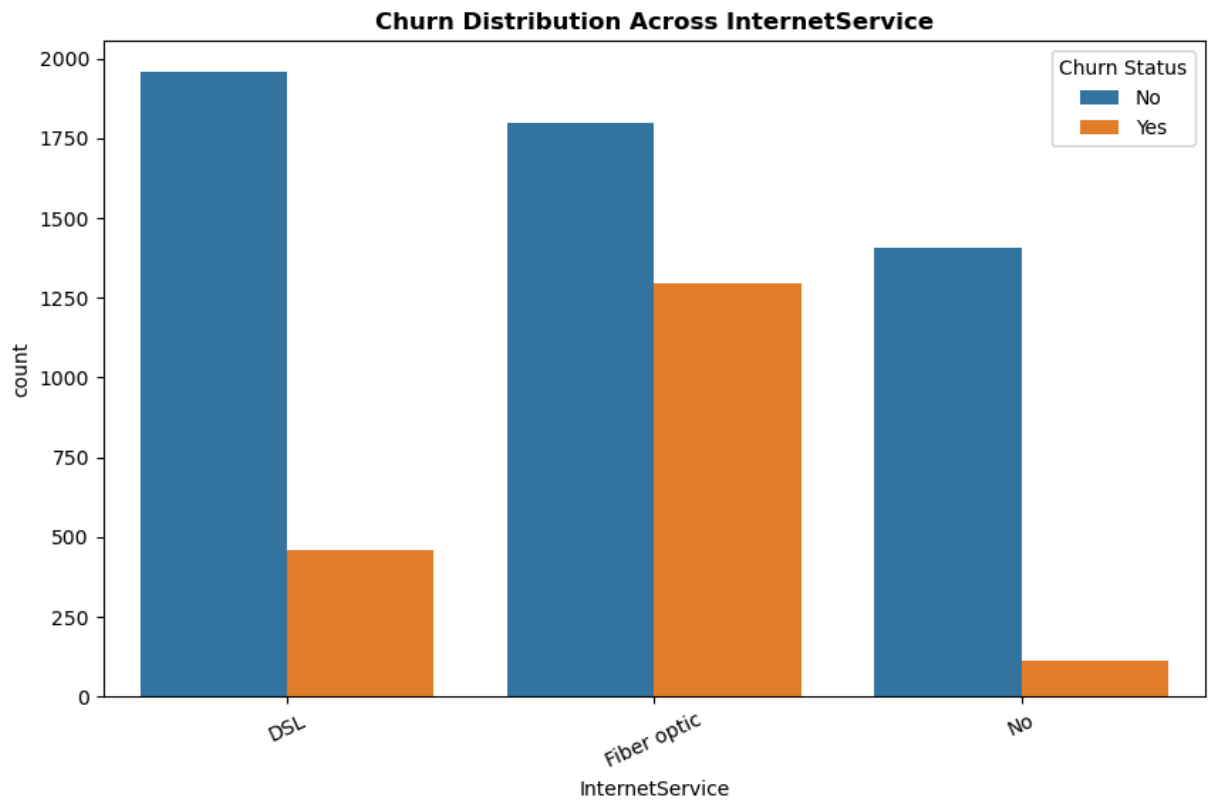
```
In [37]: import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

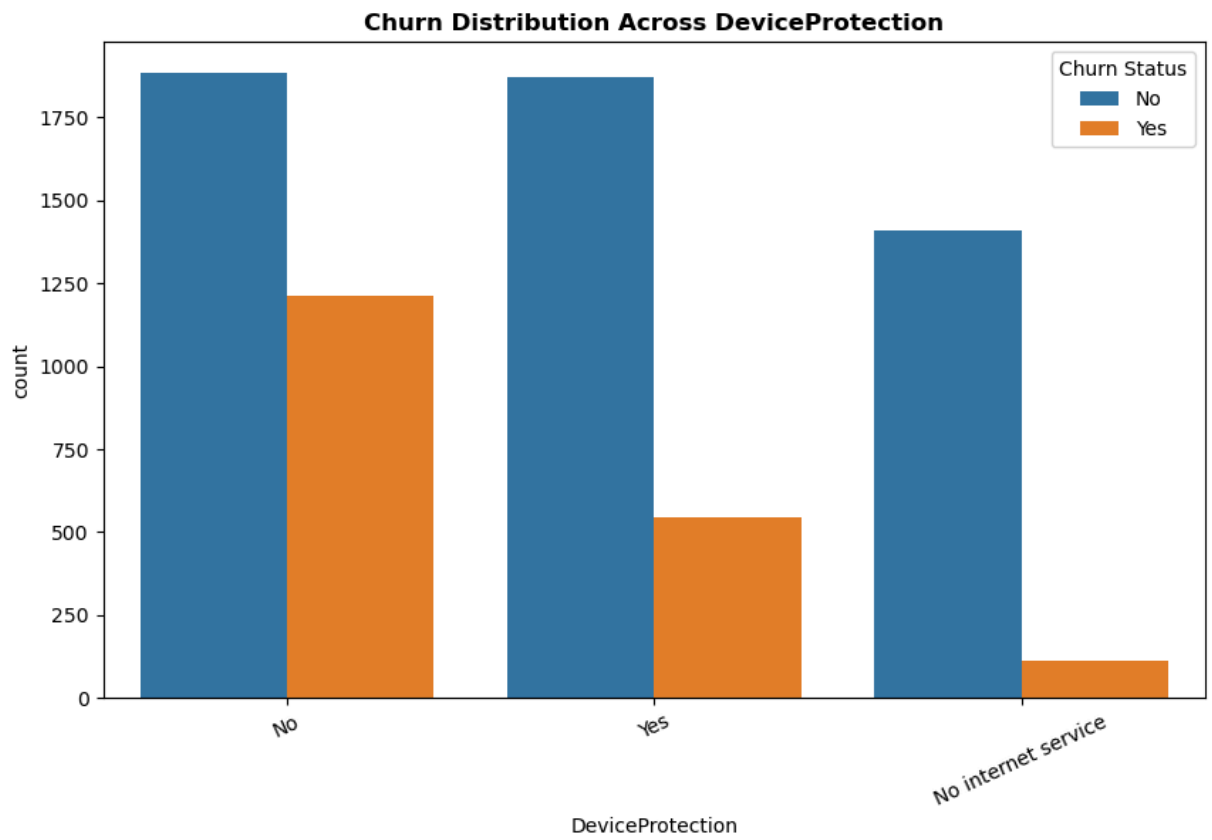
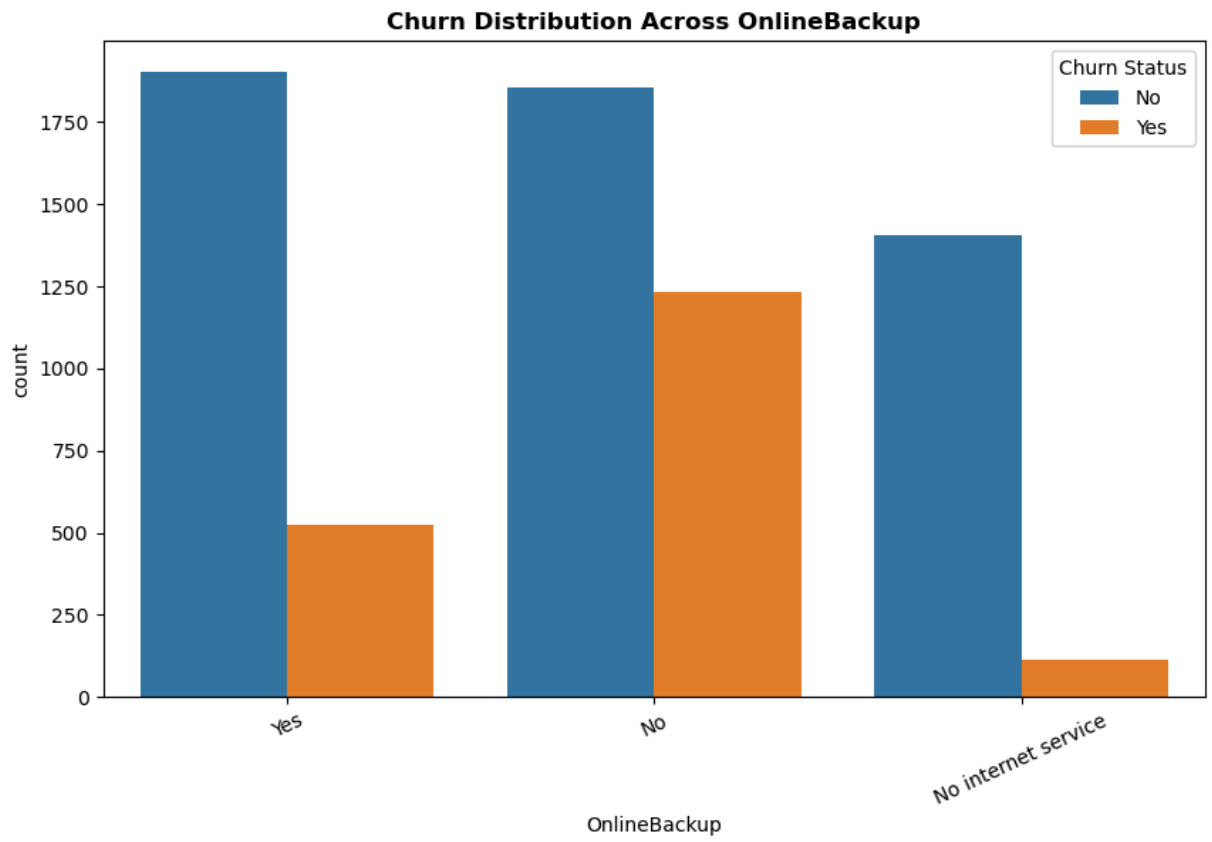
```
In [38]: # visualizing the distribution of each categorical feature with respect to churn
for idx, feature in enumerate(df.drop(columns=['Churn', 'TotalCharges', 'MonthlyCharges', 'TotalCharges'],
                                       index=[0, 1, 2, 3])):
    plt.figure(idx, figsize=(10,6))
    sns.countplot(data=df, x=feature, hue='Churn')
    plt.title(f"Churn Distribution Across {feature}", fontweight='bold')
    plt.xticks(rotation=25)
    plt.legend(title='Churn Status')
    plt.show()
    print("\n\n")
```

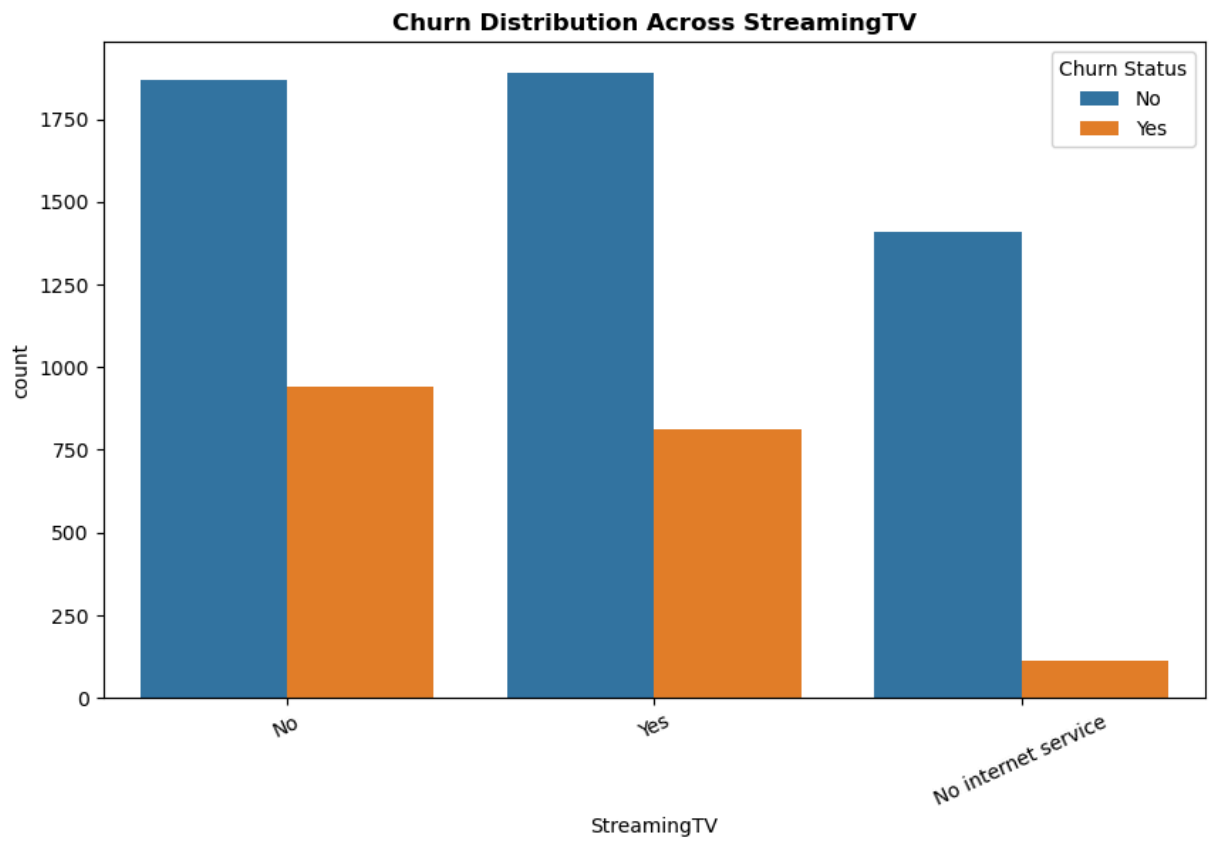
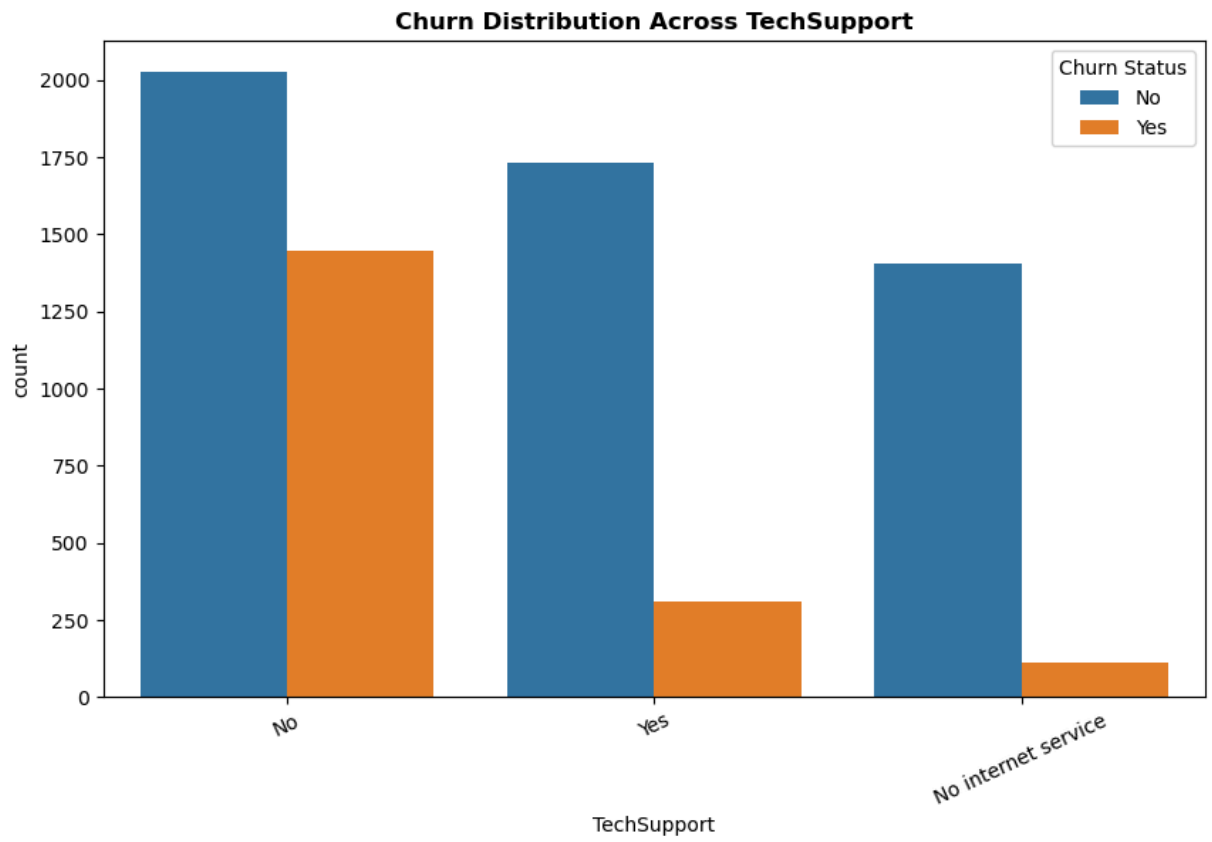


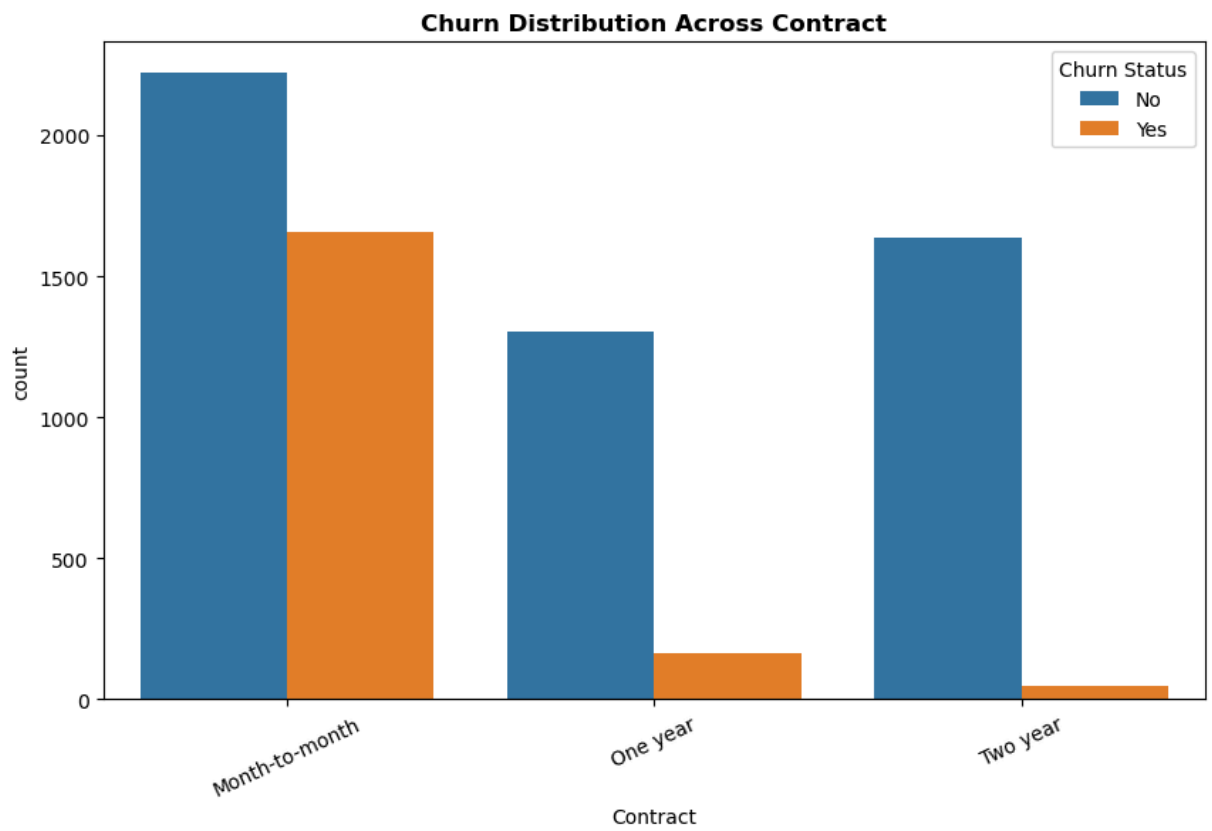
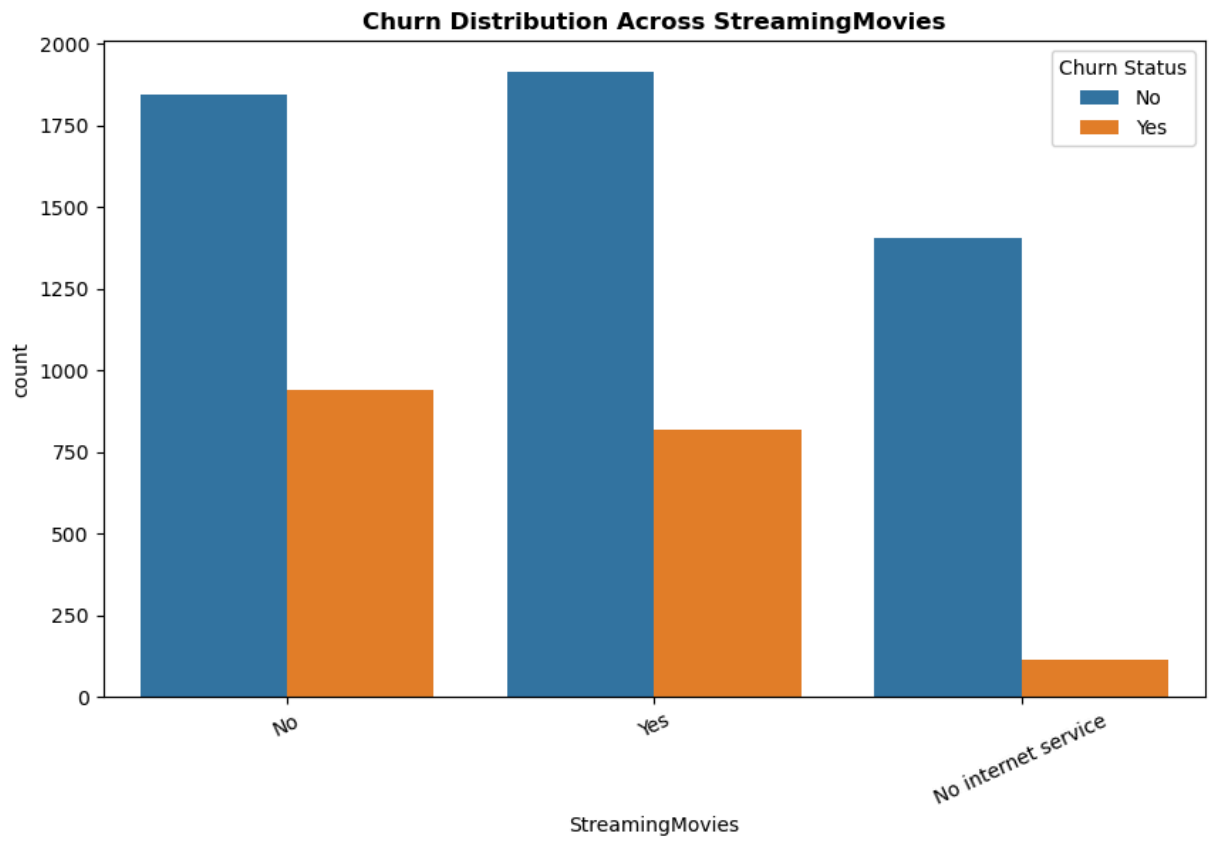




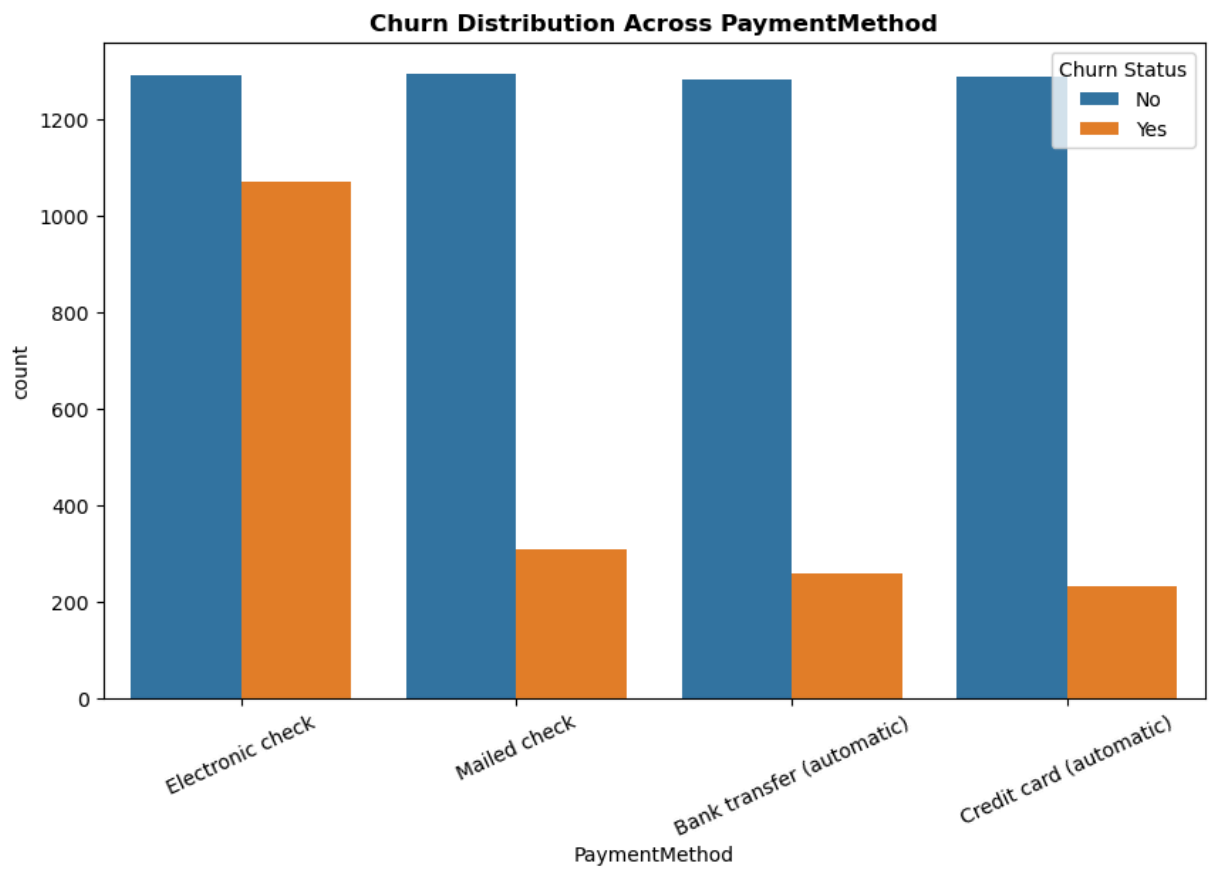
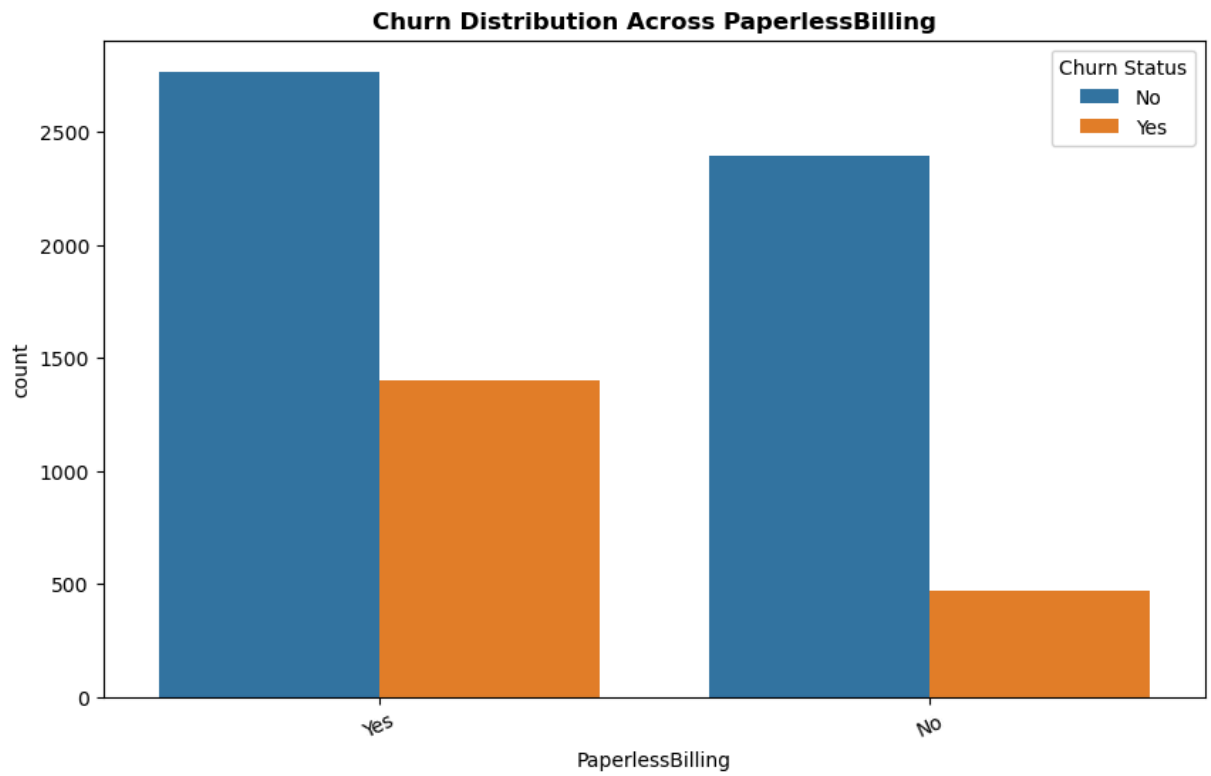


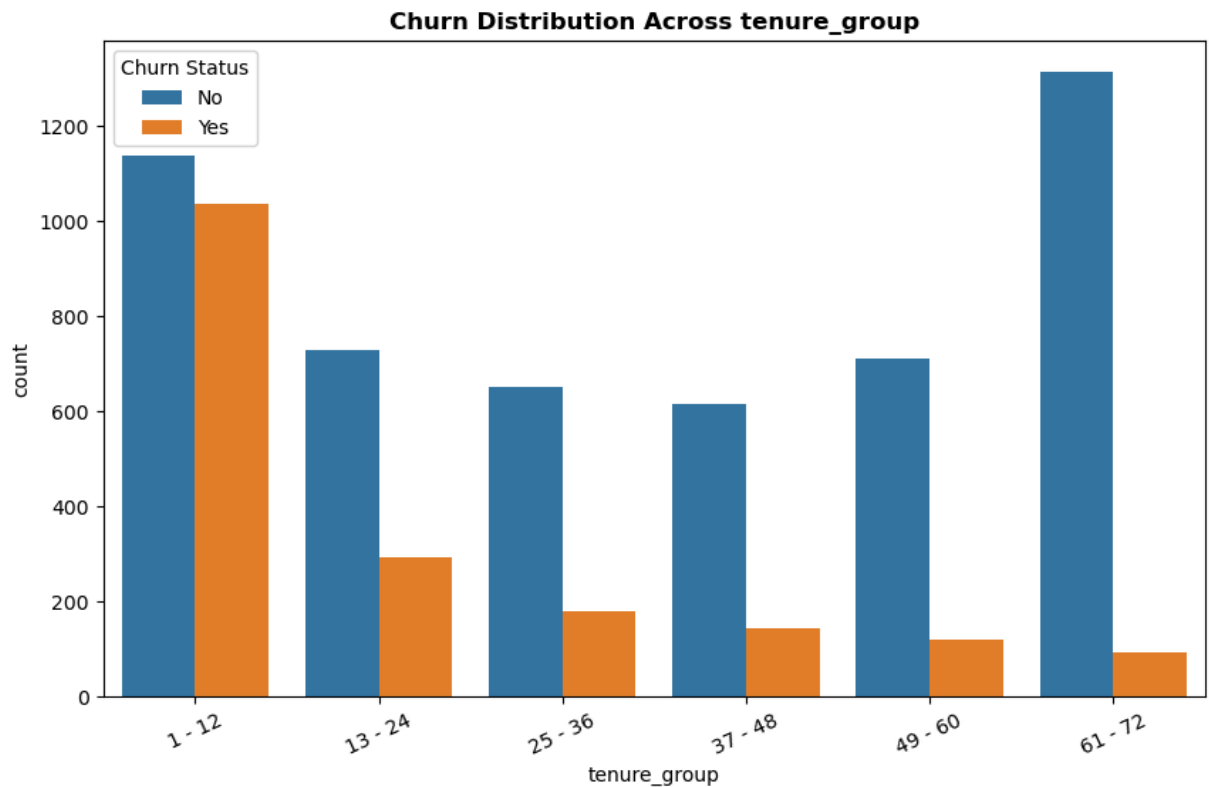












## Univariate Analysis:

- Customers who use streaming movies are less likely to churn.
- Month-to-month contracts have the highest churn.
- One-year and two-year contracts show very low churn rates.
- Churn is much higher for Fiber optic users compared to DSL users.
- Paperless billing users churn more compared to those who receive paper bills.
- Customers paying through electronic checks tend to churn the most.
- Automatic payments (bank transfer or credit card) users are more stable.
- New customers with less than a year of tenure churn the most.
- Customers with higher tenure rarely churn.

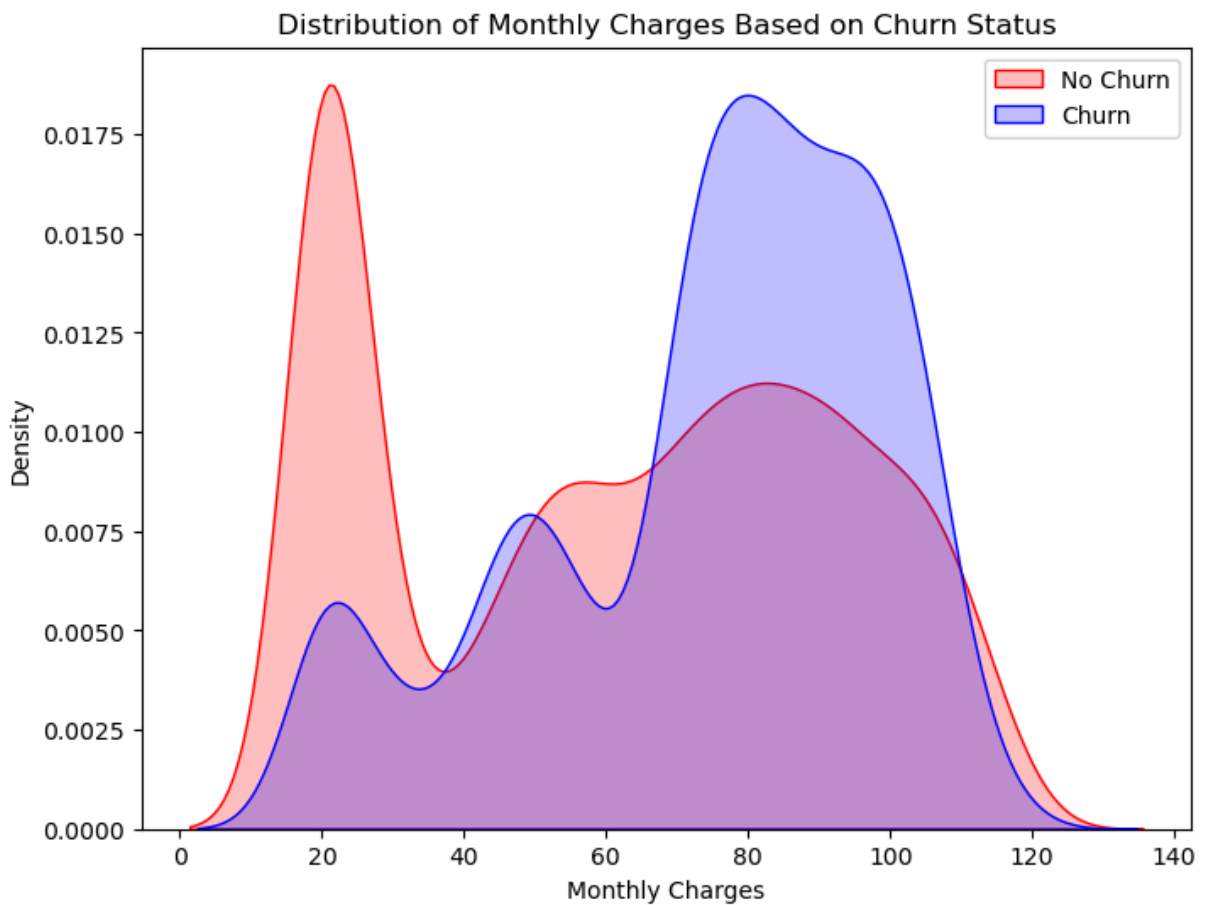
```
In [40]: df.head()
```

Out[40]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Multipl
0	Female	No	Yes	No	1	No	Nc
1	Male	No	No	No	34	Yes	
2	Male	No	No	No	2	Yes	
3	Male	No	No	No	45	No	Nc
4	Female	No	No	No	2	Yes	

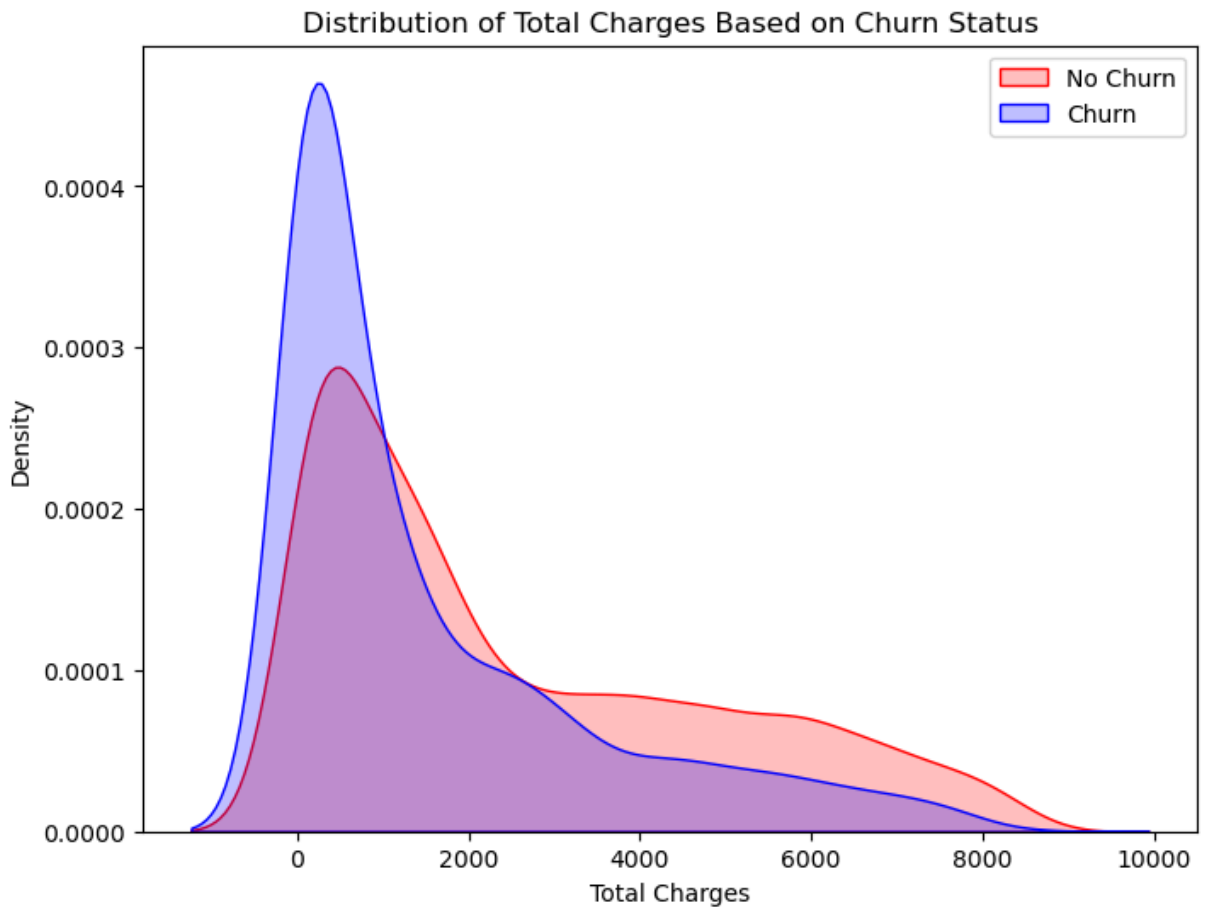
5 rows × 21 columns

```
In [41]: # visualizing the distribution of Monthly Charges for churned and non-churned
plt.figure(figsize=(8,6))
sns.kdeplot(df[df["Churn"]=="No"]["MonthlyCharges"], shade=True, color='red')
sns.kdeplot(df[df["Churn"]=="Yes"]["MonthlyCharges"], shade=True, color='blue')
plt.legend(loc='upper right')
plt.xlabel("Monthly Charges")
plt.title('Distribution of Monthly Charges Based on Churn Status')
plt.show()
```



**Insight:** Customers with higher monthly charges tend to have a higher churn rate.

```
In [43]: # visualizing the distribution of Total Charges for churned and non-churned
plt.figure(figsize=(8,6))
sns.kdeplot(df[df['Churn']=='No']['TotalCharges'], shade=True, color='red',
sns.kdeplot(df[df['Churn']=='Yes']['TotalCharges'], shade=True, color='blue'
plt.legend(loc='upper right')
plt.xlabel('Total Charges')
plt.title("Distribution of Total Charges Based on Churn Status")
plt.show()
```



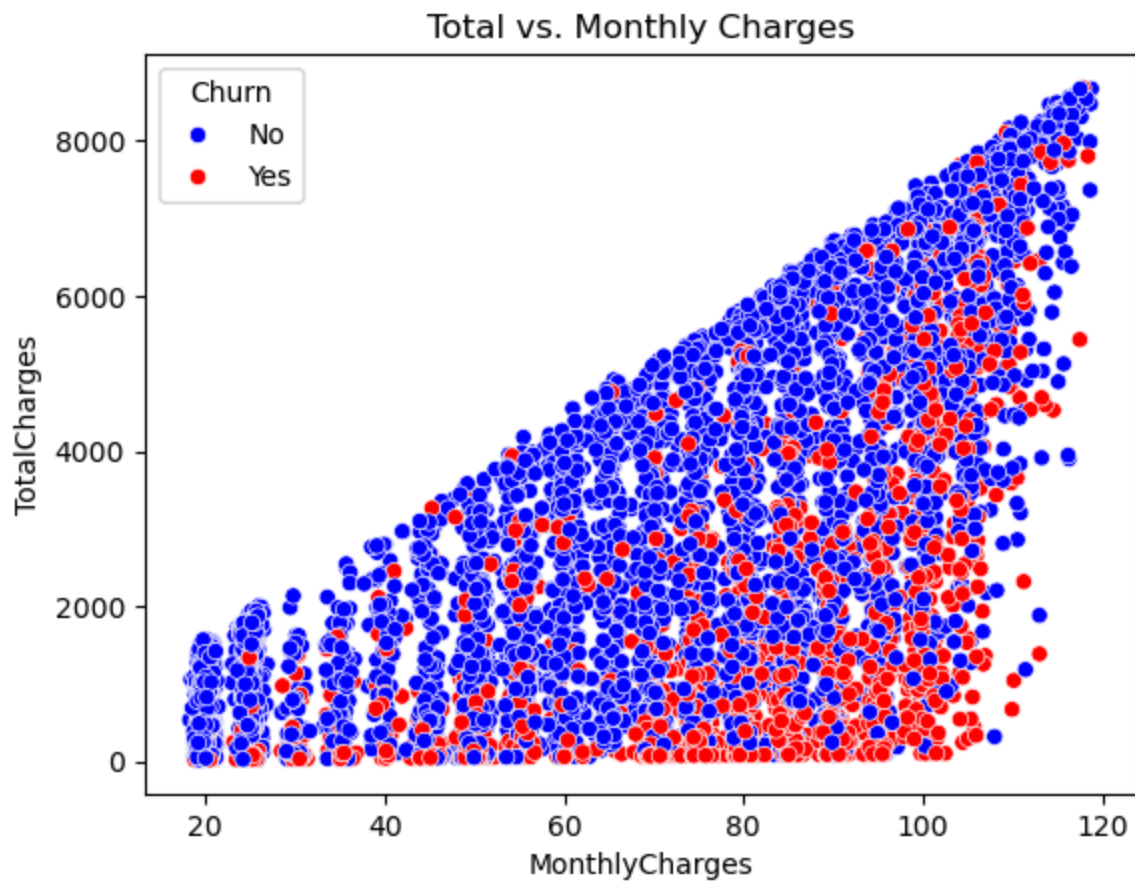
### Key Insight:

Churn tends to be higher when **Total Charges** are lower.

However, when analyzing **Tenure**, **Monthly Charges**, and **Total Charges** together, a clearer pattern emerges—**higher Monthly Charges with shorter Tenure lead to lower Total Charges**. This suggests that customers with **high Monthly Charges, low Tenure, and low Total Charges** are more likely to churn.

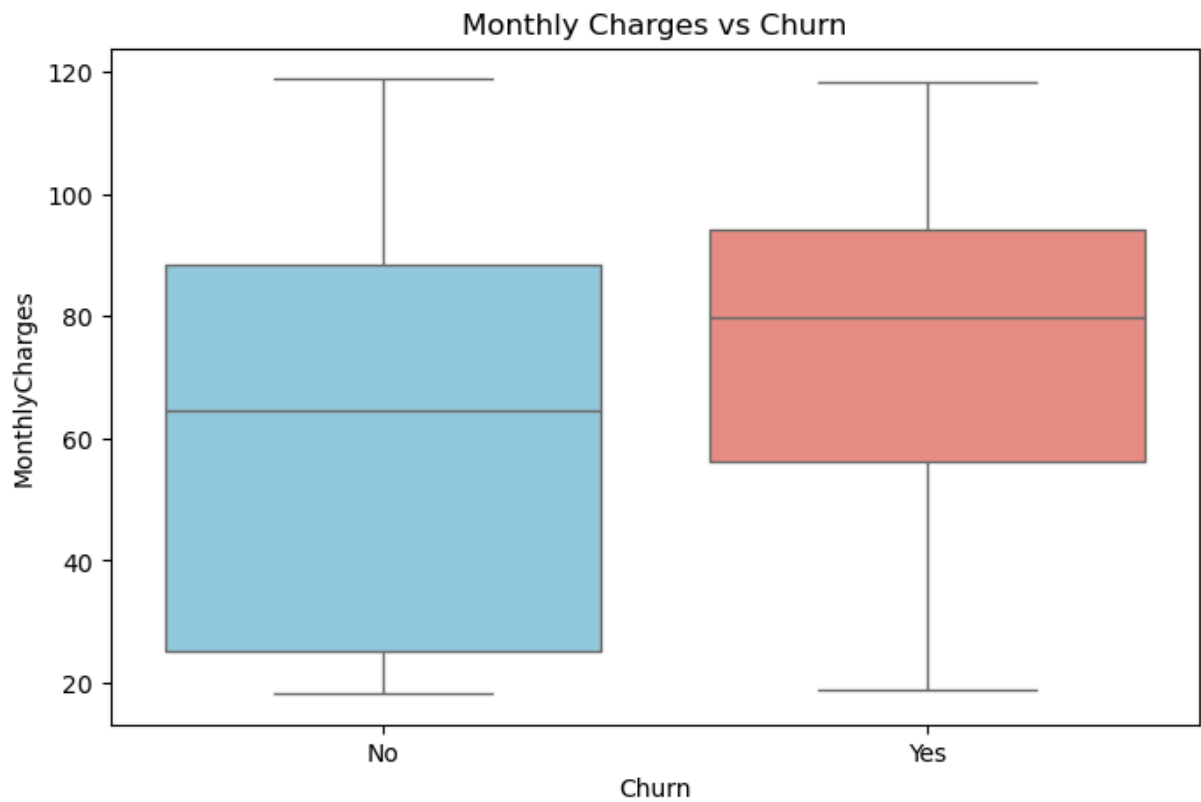
## 5 Bivariate Analysis

```
In [46]: # Relation between Monthly Charges and Total Charges
sns.scatterplot(df, x='MonthlyCharges', y='TotalCharges', hue='Churn', palette=
plt.title("Total vs. Monthly Charges")
plt.show()
```

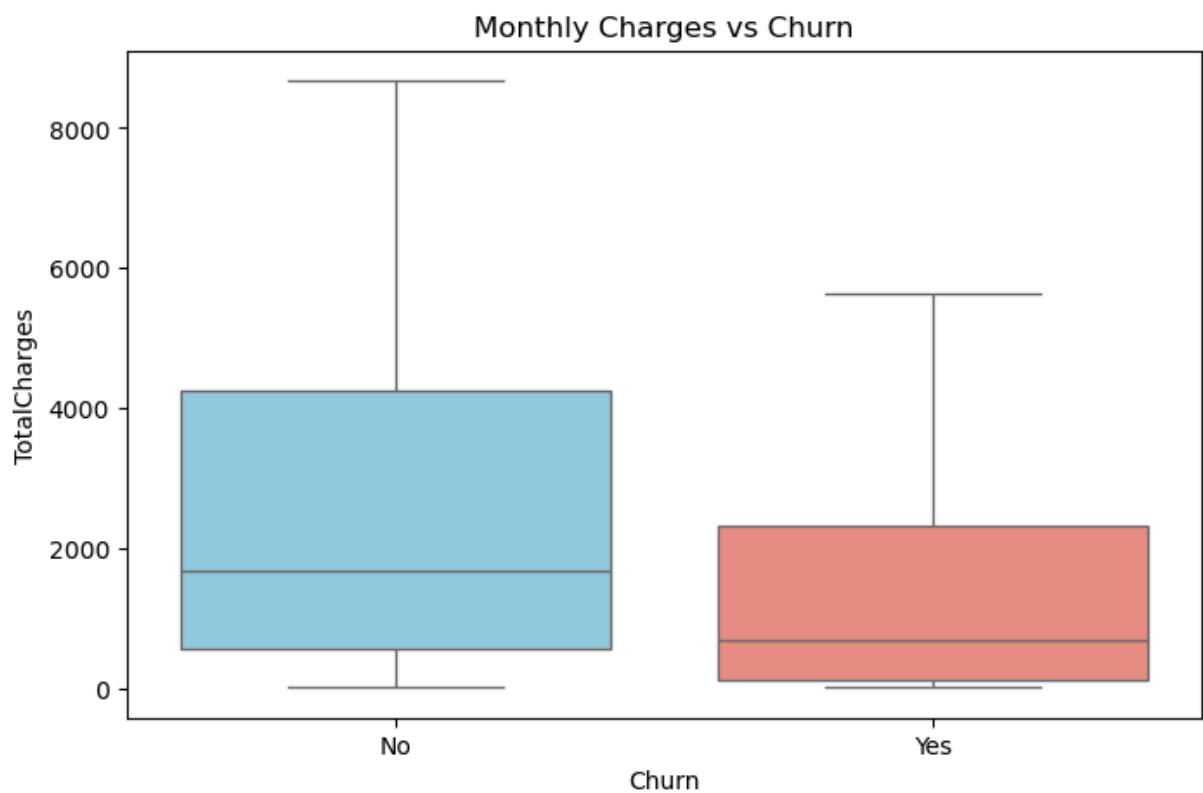


**Total Charges** increase as **Monthly Charges** increase.

```
In [48]: plt.figure(figsize=(8,5))
sns.boxplot(df,x="Churn", y='MonthlyCharges', palette=["skyblue", "salmon"])
plt.title('Monthly Charges vs Churn')
plt.show()
```



```
In [49]: plt.figure(figsize=(8,5))
sns.boxplot(df,x="Churn", y='TotalCharges', showfliers=False, palette=["skyb
plt.title('Monthly Charges vs Churn')
plt.show()
```



```
In [50]: # splitting data based on churn status
```

```
churn_yes = df[df["Churn"]=="Yes"]  
churn_no = df[df["Churn"]=="No"]
```

```
In [51]: # custom plotting function for categorical variables
```

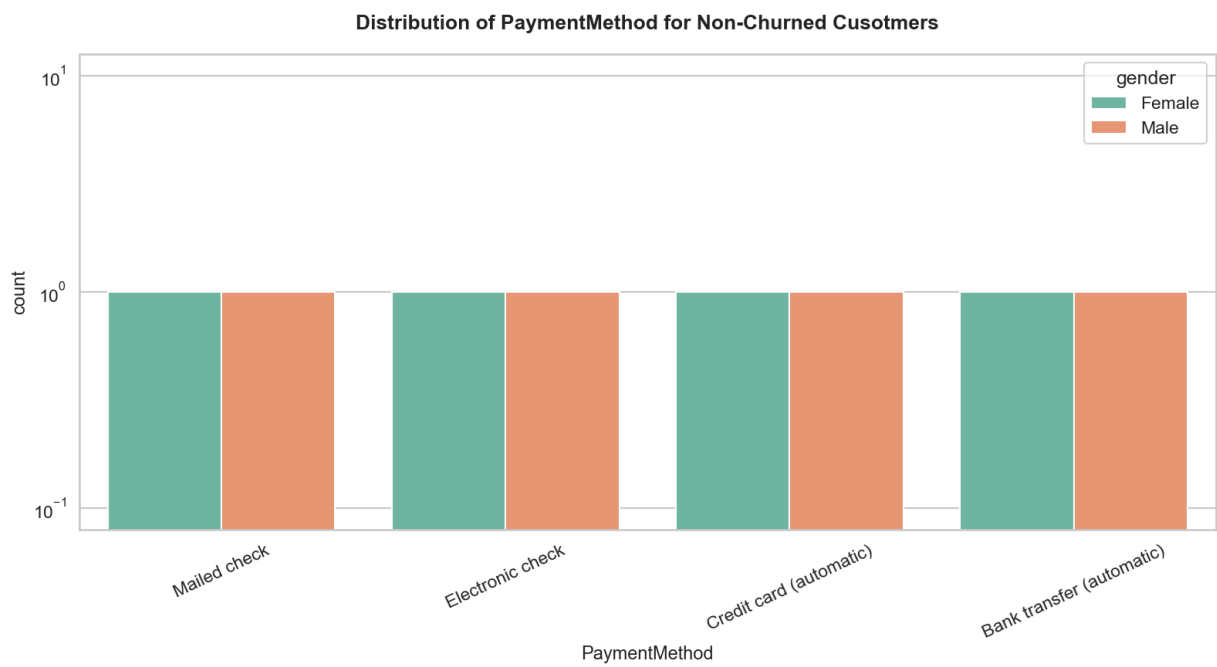
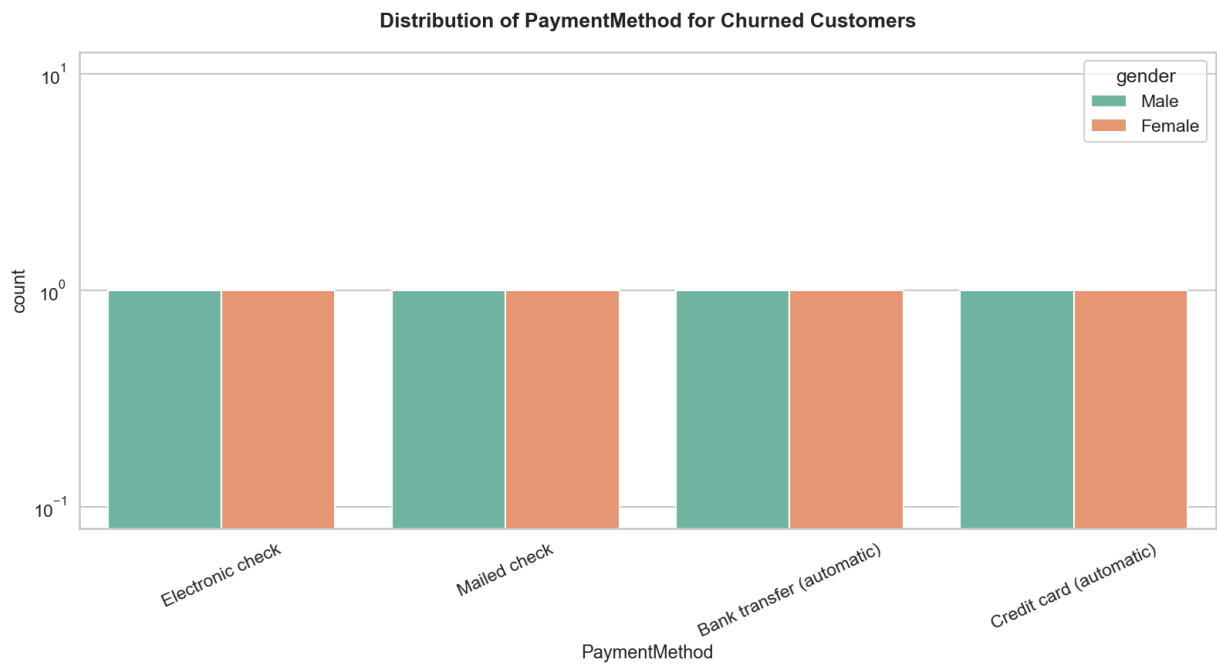
```
def custom_plot(data, feature, title, hue=None):  
    sns.set_style('whitegrid')  
    sns.set_context('talk')  
    plt.rcParams.update({'axes.labelsize':17, 'axes.titlesize':19, 'axes.tit  
  
    feature_unique_count = data[feature].nunique()  
    hue_unique_count = data[hue].nunique() if hue else 0  
    width = feature_unique_count + 7 + 4 * hue_unique_count  
  
    fig, ax = plt.subplots(figsize=(width, 8))  
  
    plt.xticks(rotation=25)  
    plt.yscale('log')  
    plt.title(title, fontweight='bold')  
  
    sns.countplot(data=data, x=feature, hue=hue, order=data[feature].value_c  
    plt.show()
```

```
In [52]: df.columns.values
```

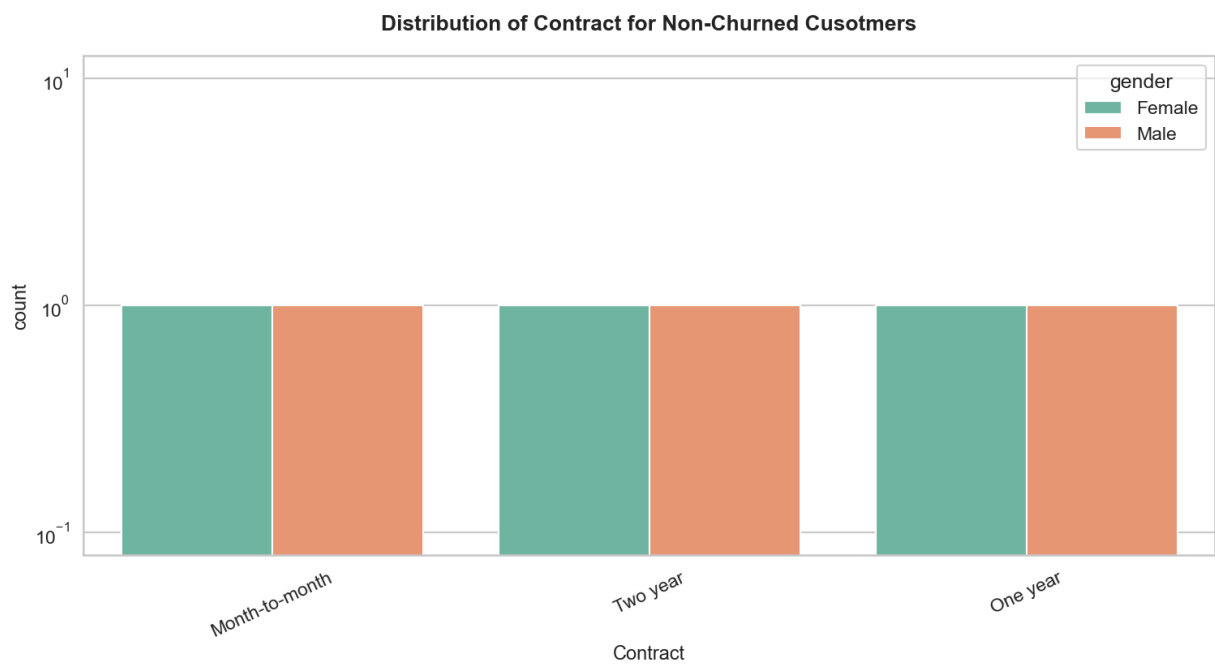
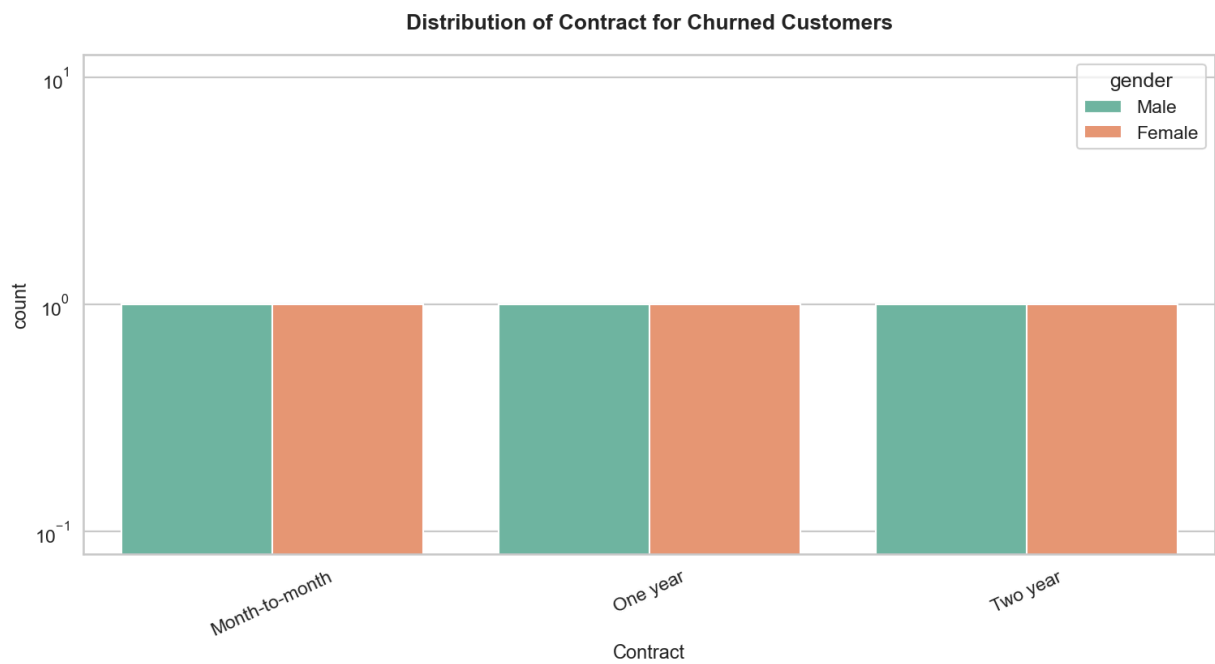
```
Out[52]: array(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',  
               'PhoneService', 'MultipleLines', 'InternetService',  
               'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',  
               'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',  
               'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',  
               'TotalCharges', 'Churn', 'tenure_group'], dtype=object)
```

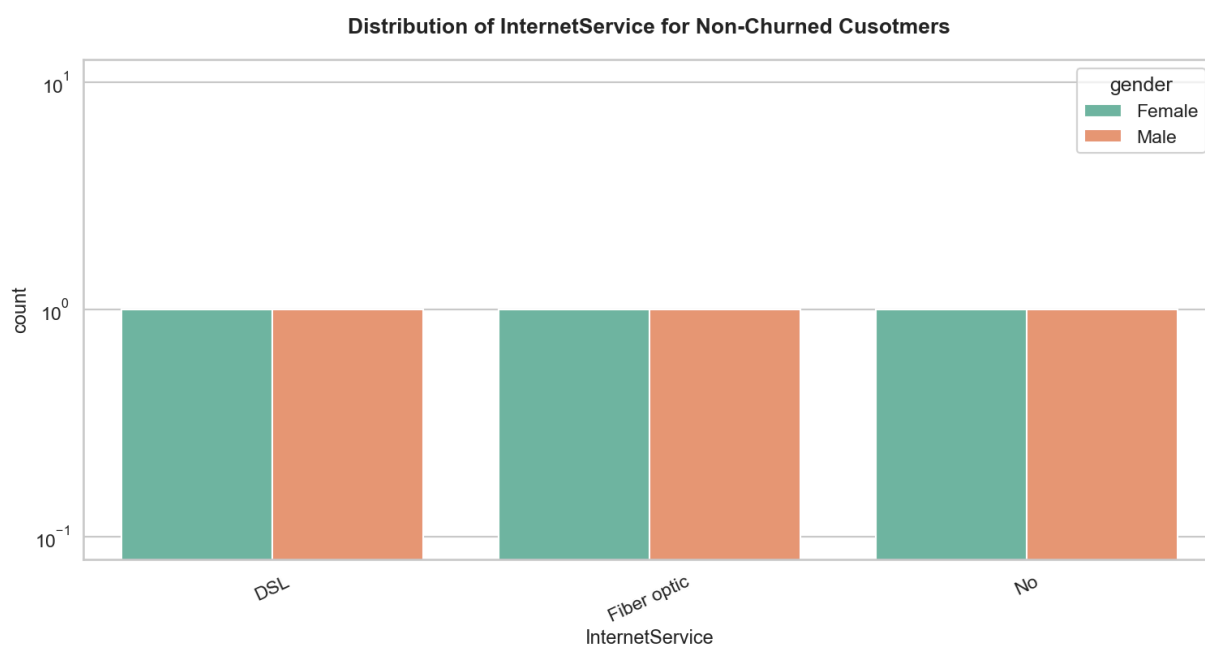
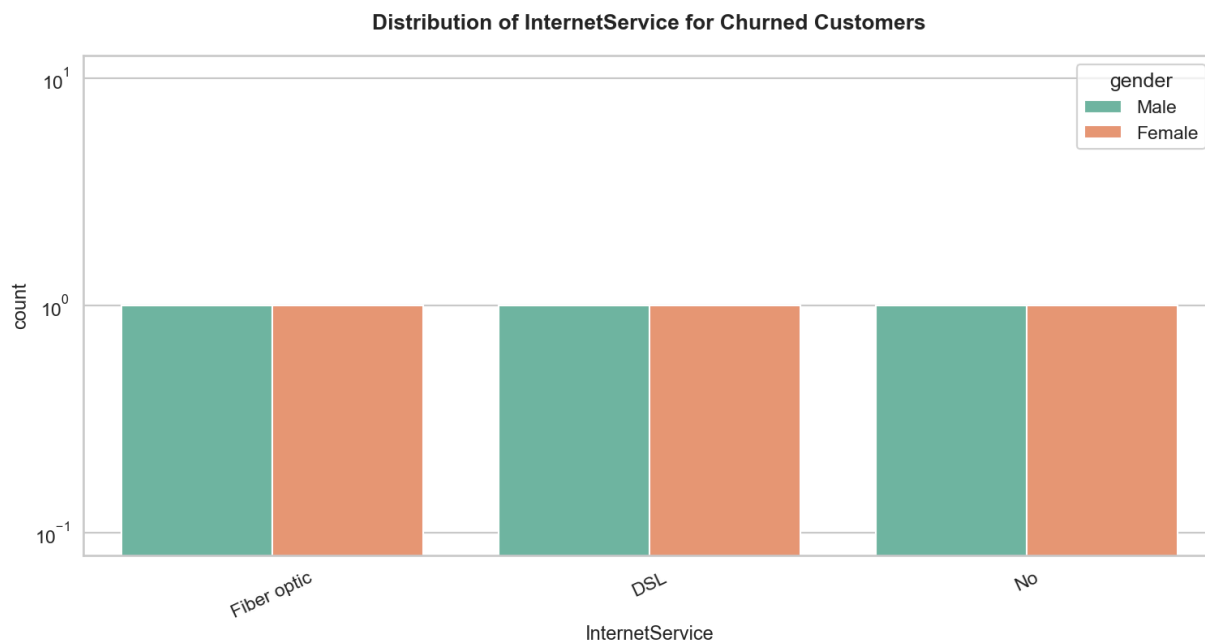
```
In [53]: cat_features = ['PaymentMethod', 'Contract', 'InternetService', 'Partner', ' '
```

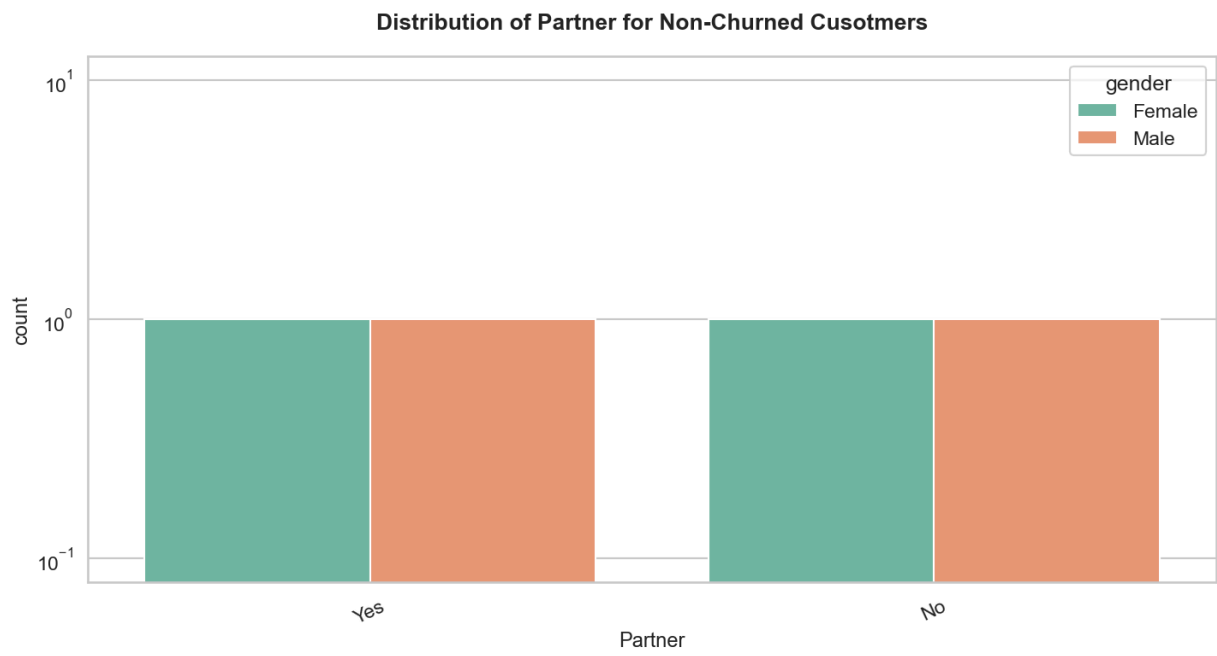
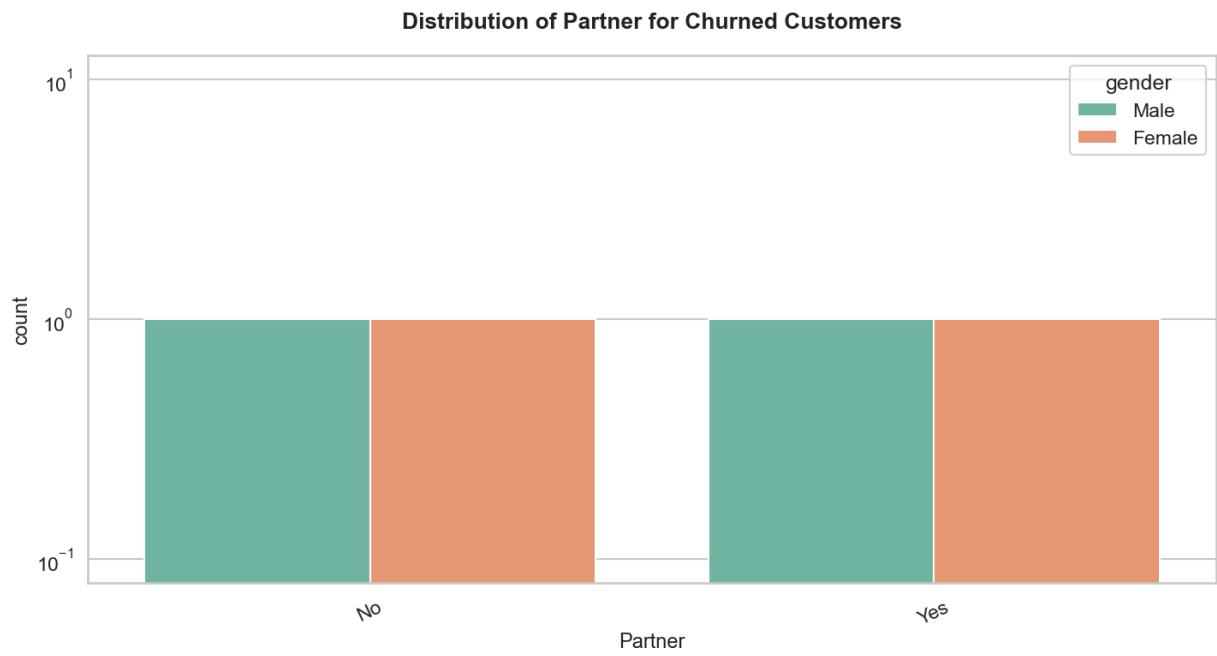
```
for feature in cat_features:  
    # plot for churned customers  
    custom_plot(churn_yes, feature=feature, title=f'Distribution of {feature}  
    print("\n")  
    # plot for non churned customers  
    custom_plot(churn_no, feature=feature, title=f'Distribution of {feature}  
    print("\n")
```

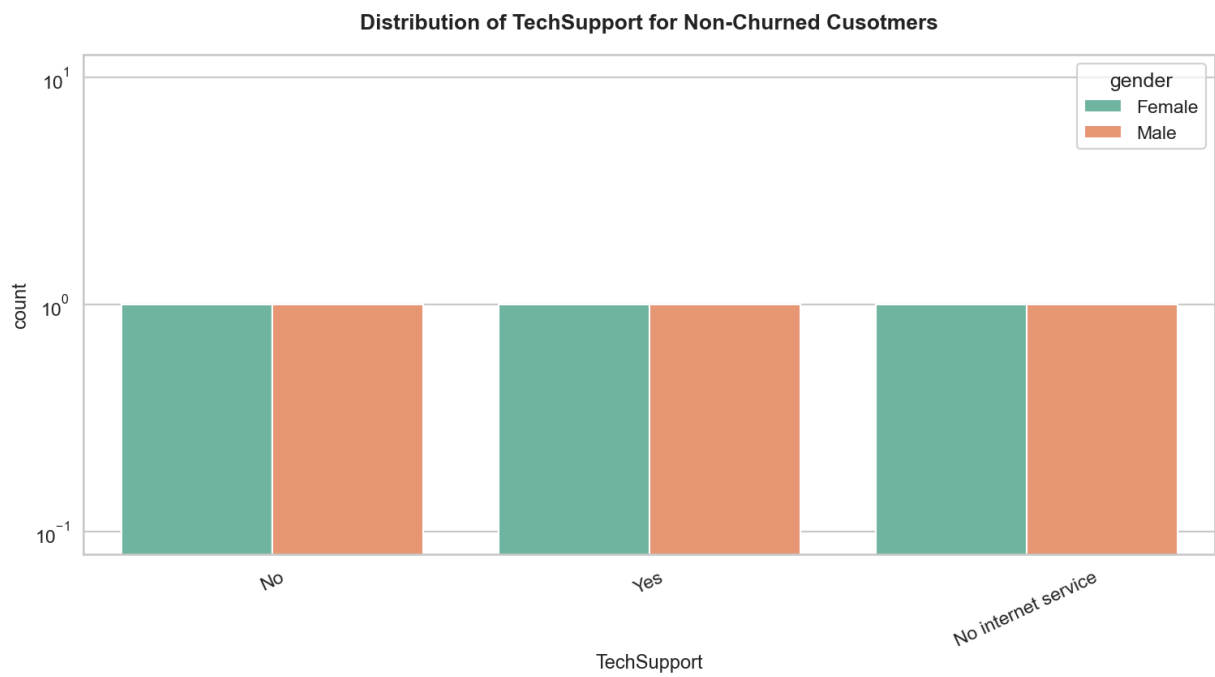
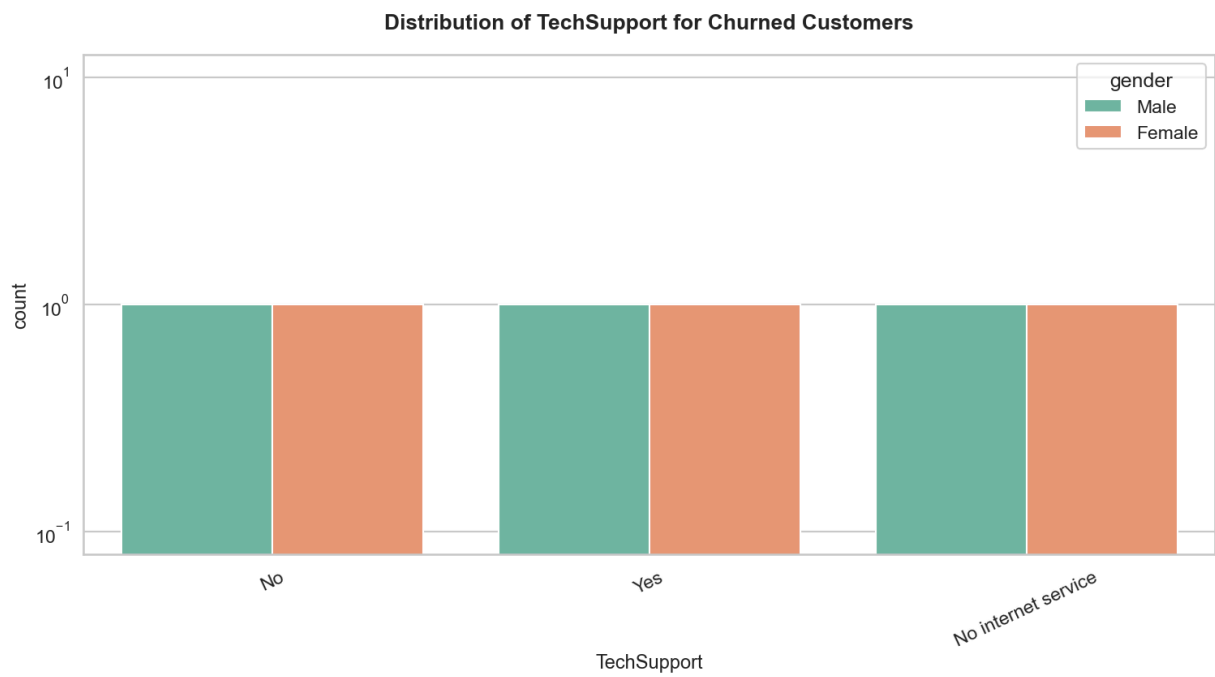


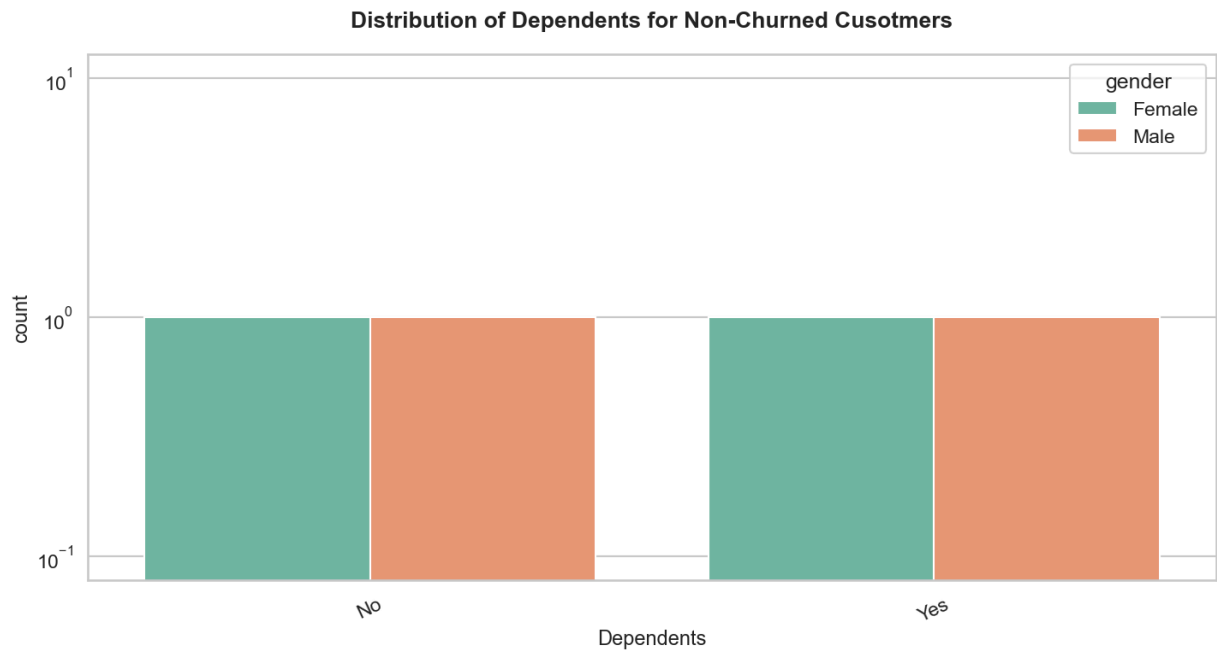
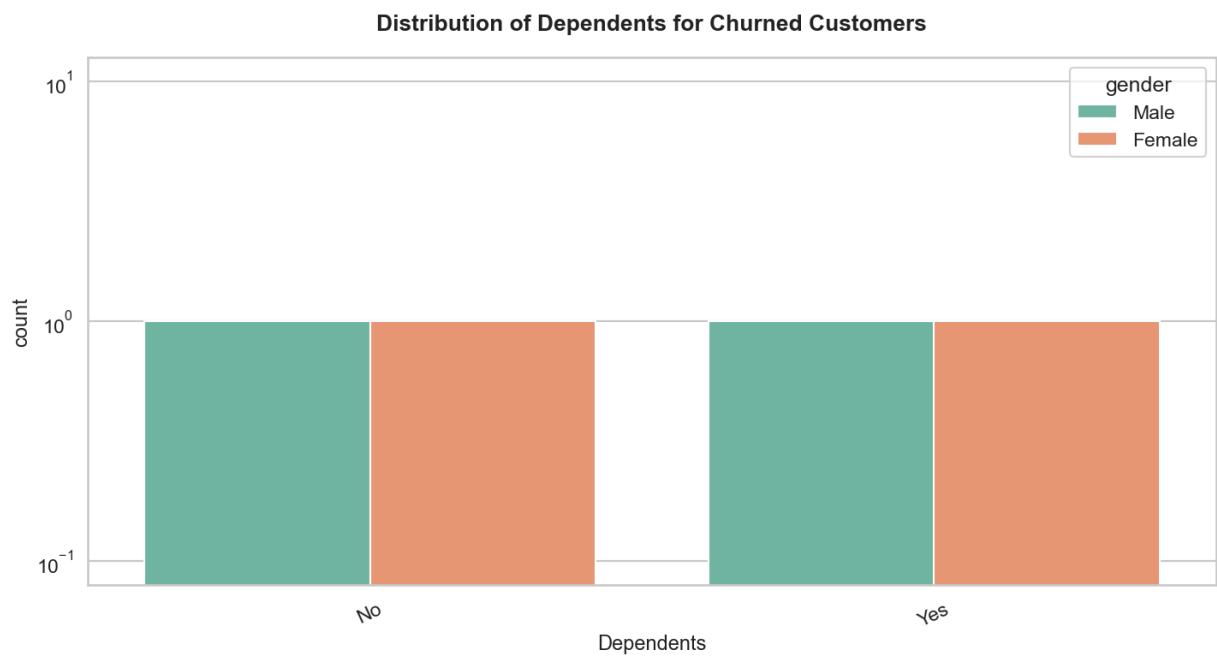




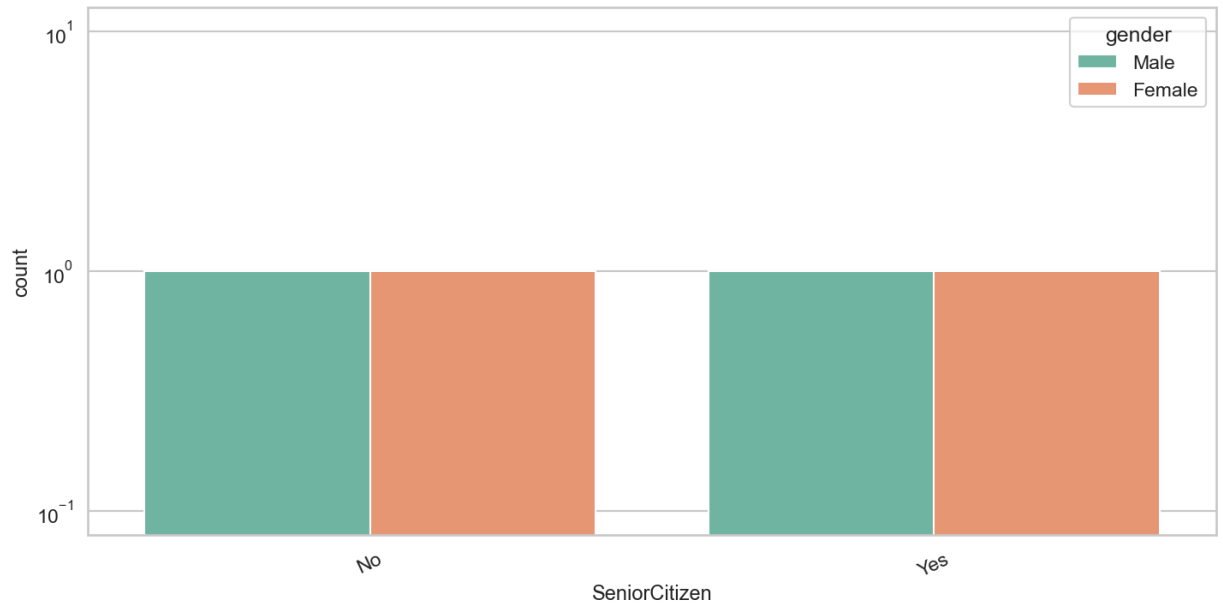




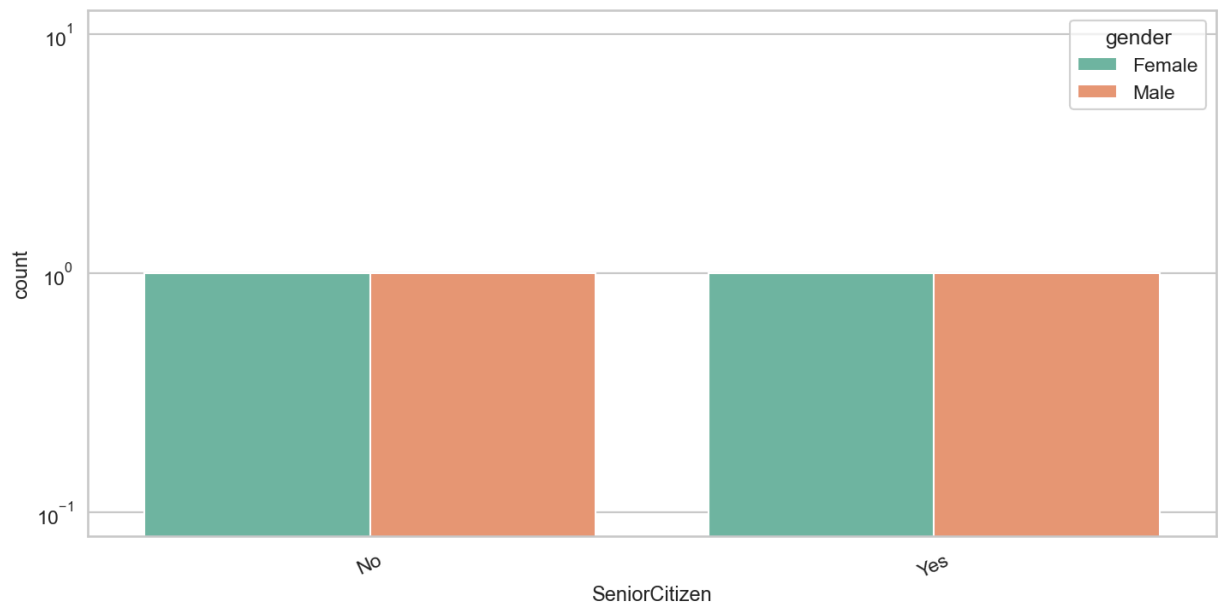


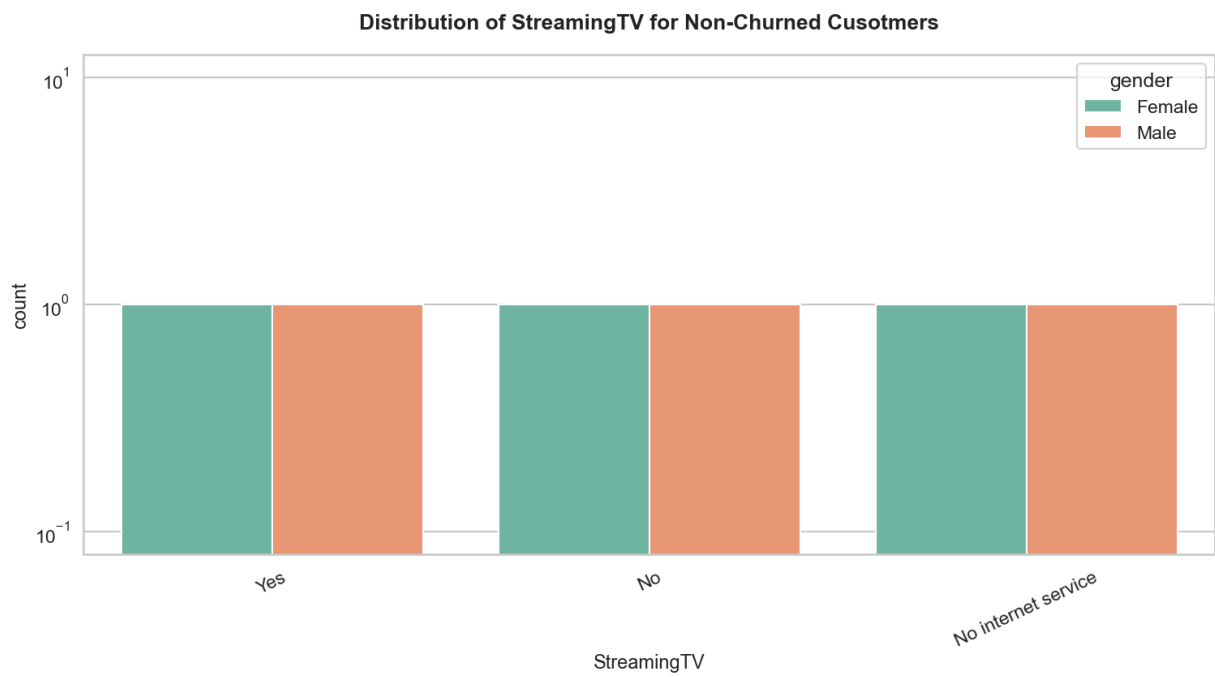
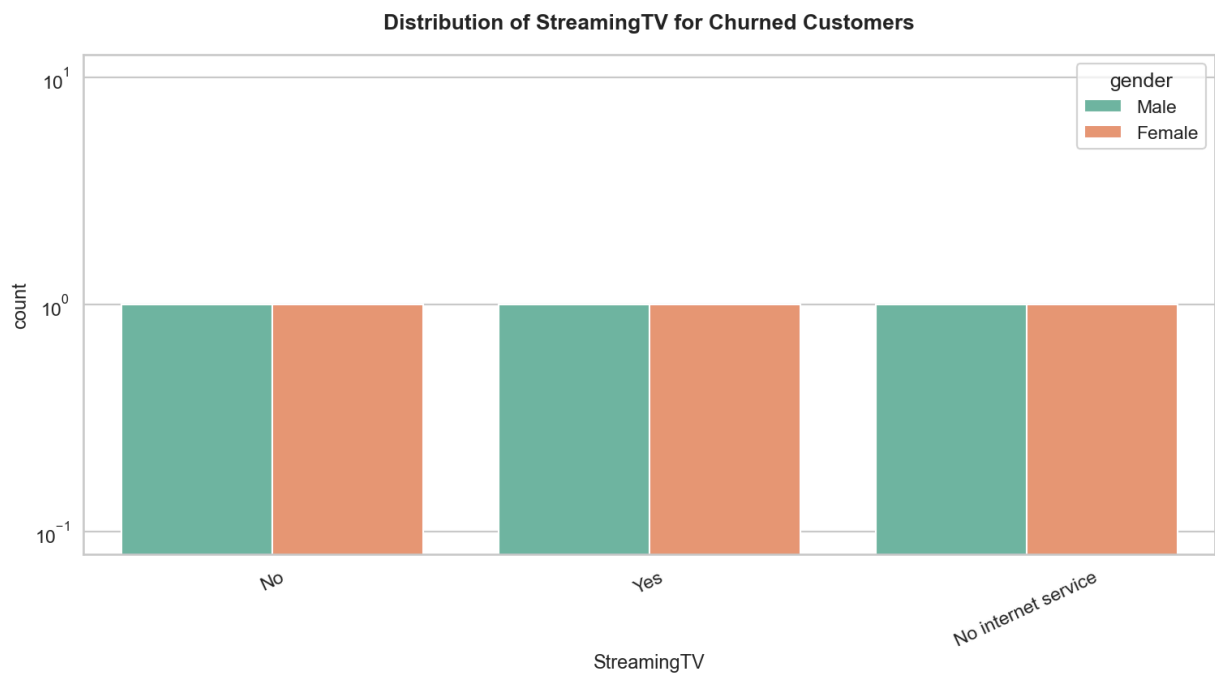


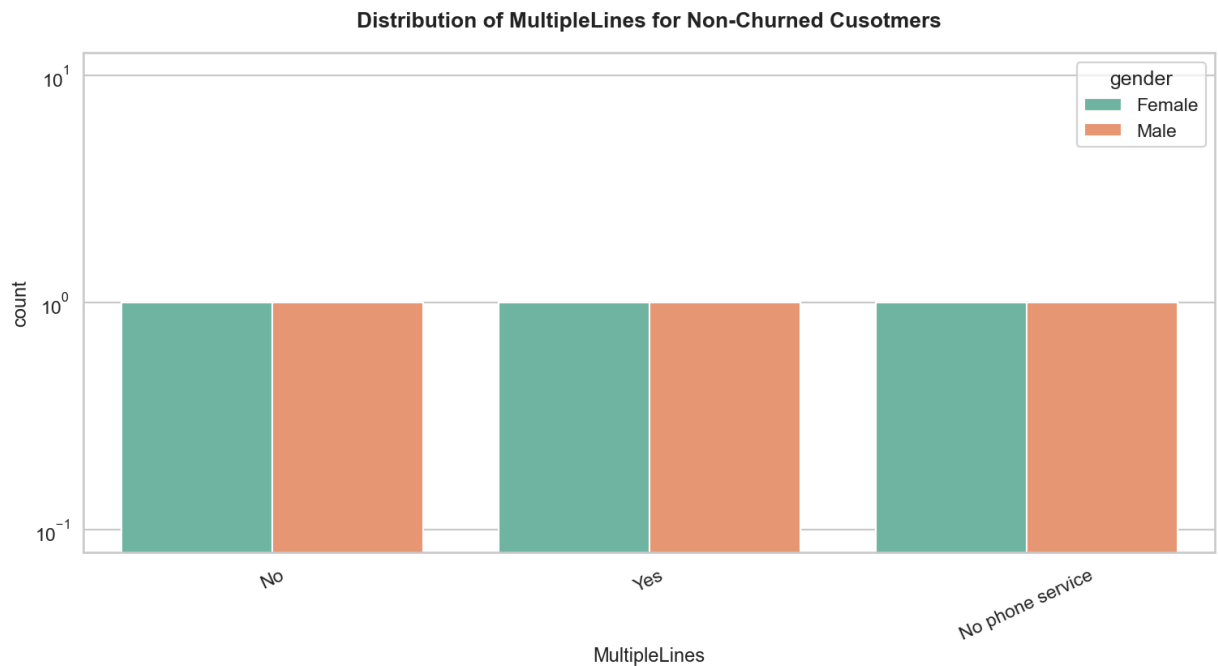
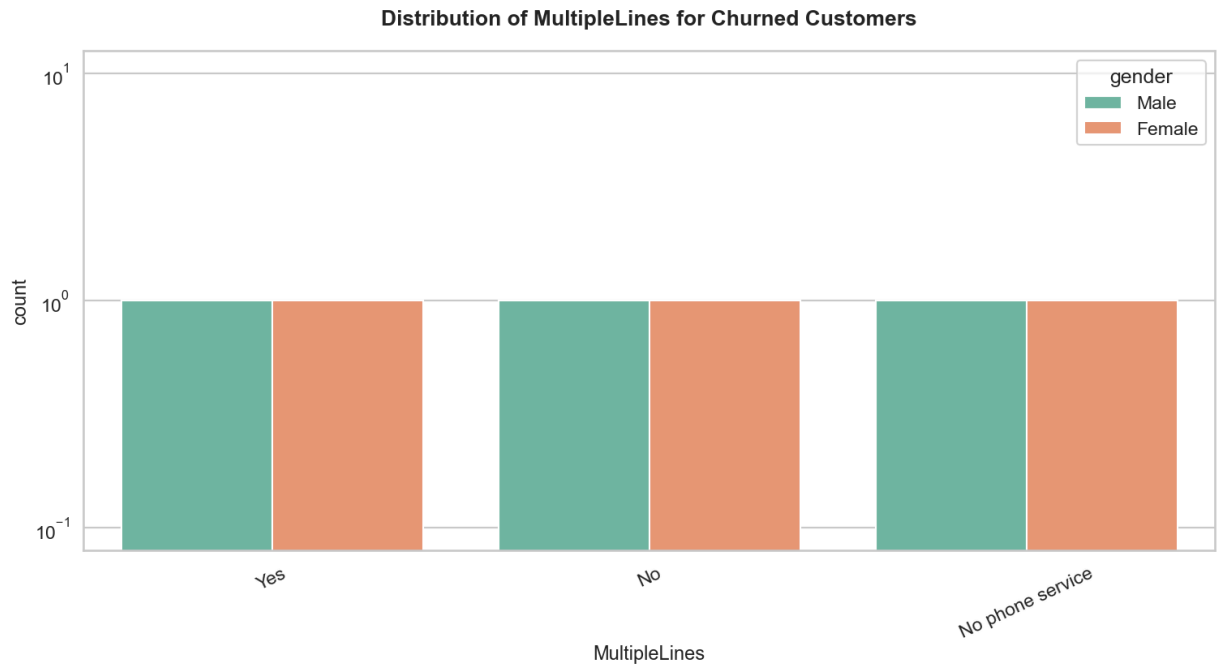
Distribution of SeniorCitizen for Churned Customers



Distribution of SeniorCitizen for Non-Churned Customers







## Bivariate Analysis:

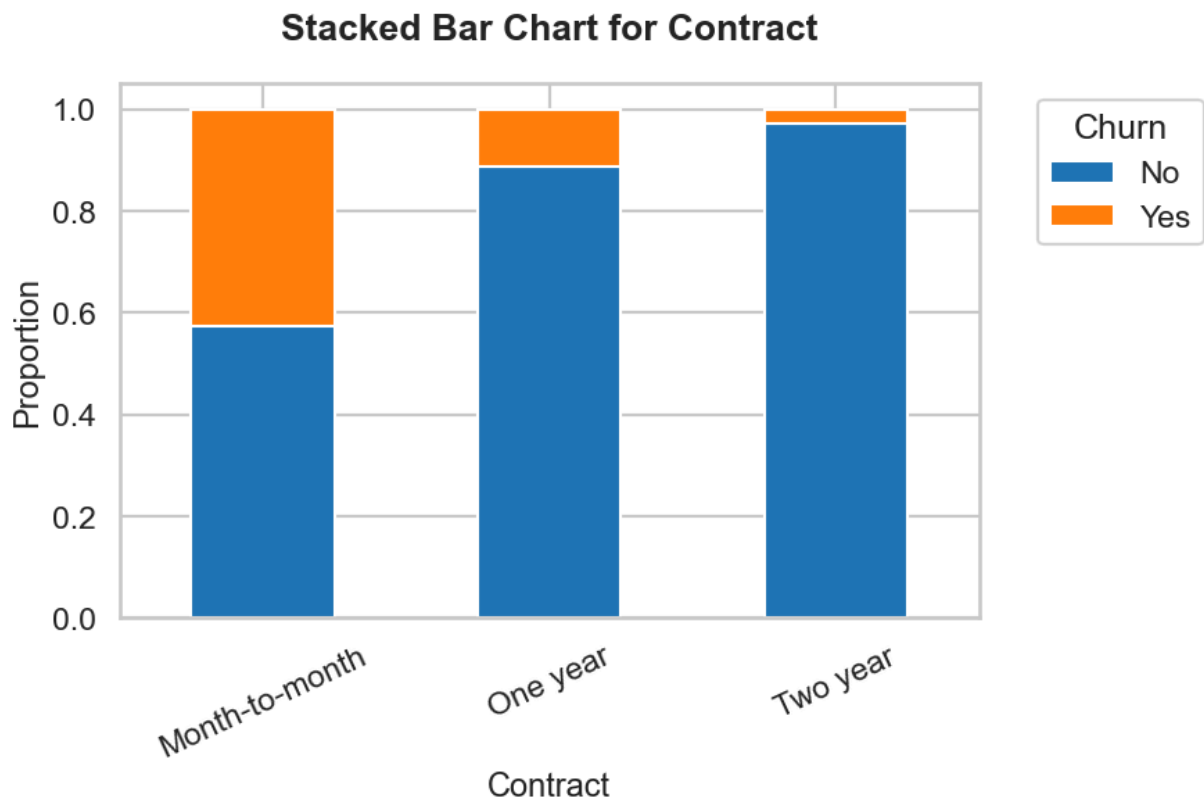
- Both males and females have similar churn patterns.
- Churned customers are almost equally split between having Streaming TV and not having it.
- Gender doesn't make much difference in churn behavior for Streaming TV users.
- Fiber optics Internet Service users are more churners.
- DSL InternetService users are less churners.
- Non-churned customers are slightly more likely to have Streaming TV.



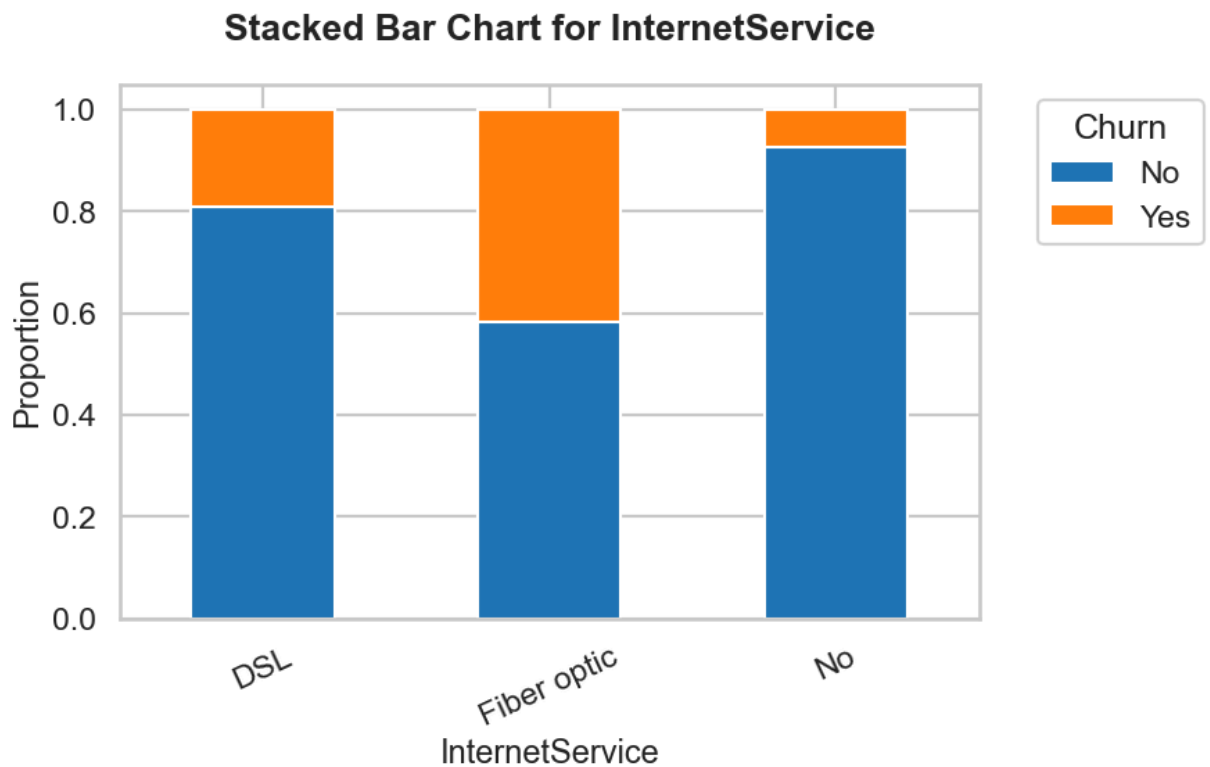
- Churned customers are also almost equally split between having multiple lines and not.
- A few churned customers don't even have phone service.
- Non-churned customers slightly lean towards not having multiple lines.
- Gender impact is minimal across all categories.

```
In [55]: def plot_stacked_bar(feature):
crosstab = pd.crosstab(df[feature], df['Churn'])
crosstab_norm = crosstab.div(crosstab.sum(axis=1), axis=0)
crosstab_norm.plot(kind='bar', stacked=True, figsize=(8,5))
plt.title(f"Stacked Bar Chart for {feature}", fontweight='bold')
plt.ylabel('Proportion')
plt.legend(title='Churn', bbox_to_anchor=(1.05, 1))
plt.xticks(rotation=25)
plt.show()
```

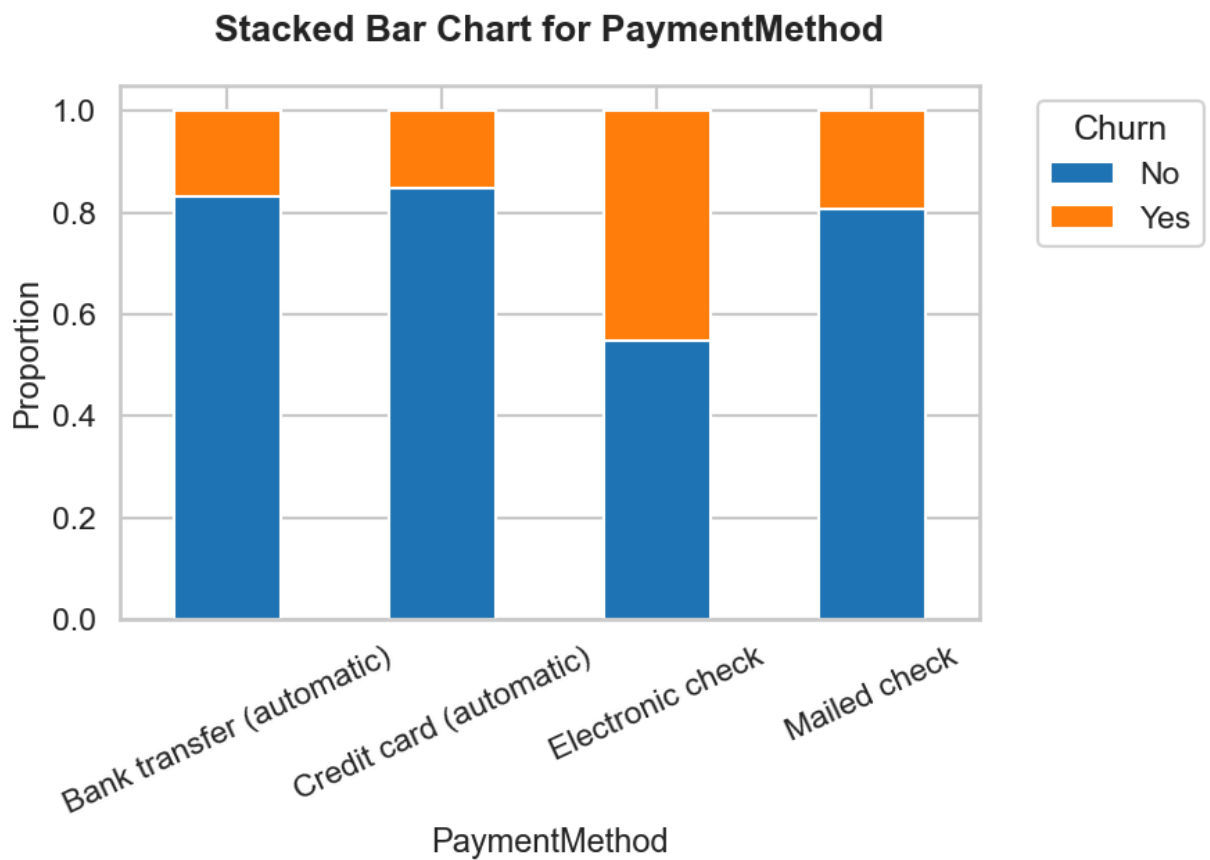
```
In [56]: plot_stacked_bar('Contract')
```



```
In [57]: plot_stacked_bar('InternetService')
```

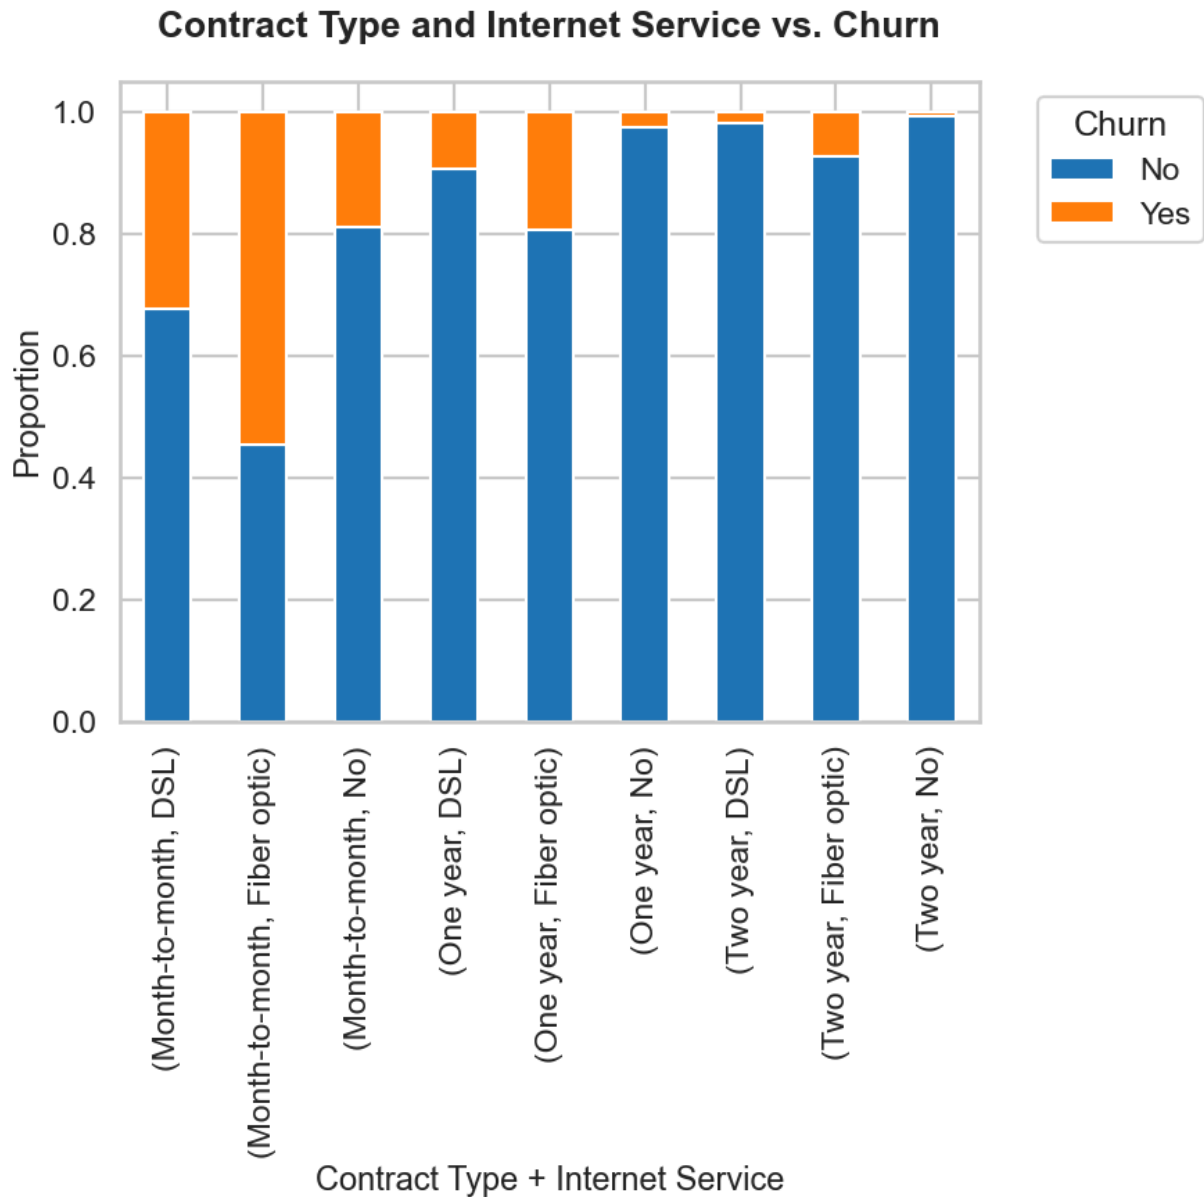


```
In [58]: plot_stacked_bar('PaymentMethod')
```



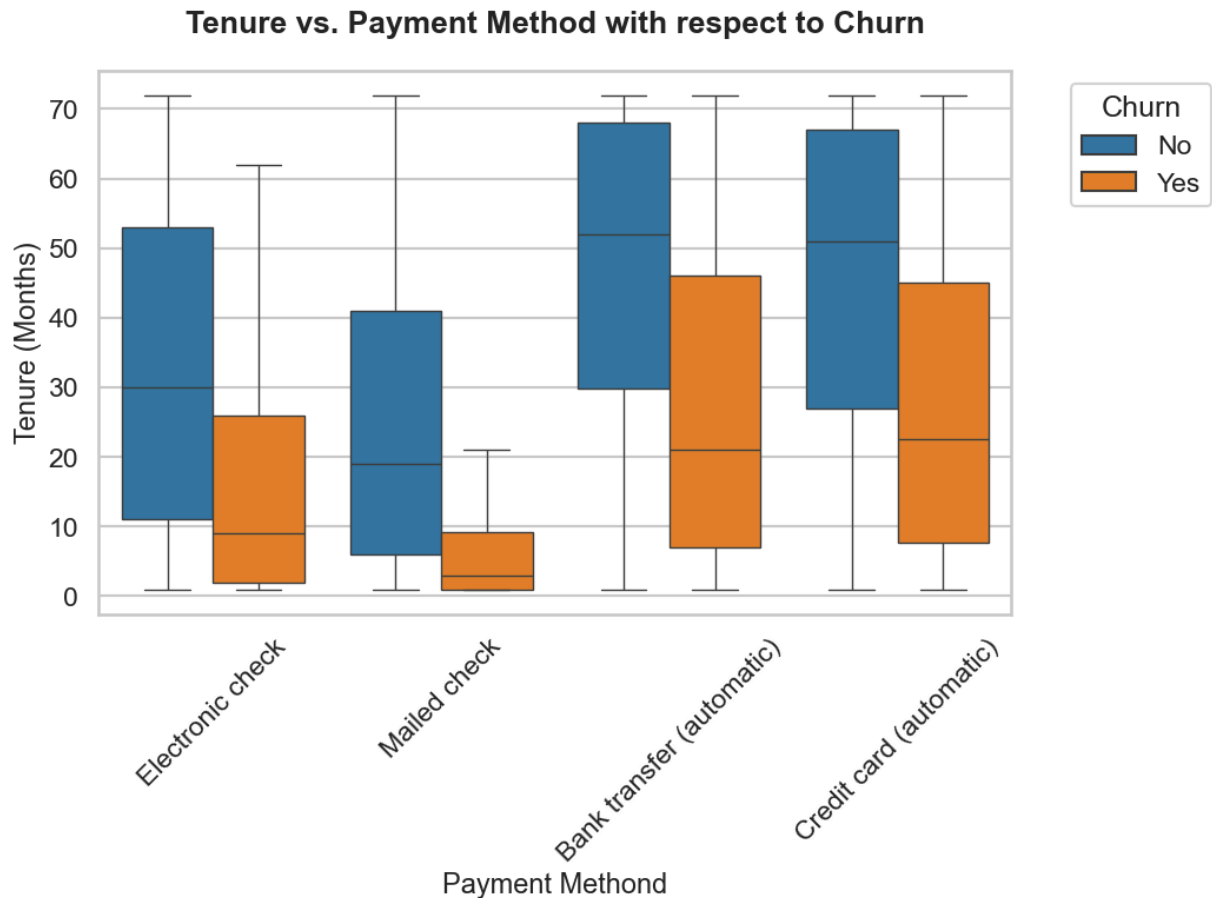
## 6 Multivariate Analysis

```
In [60]: # contract + internet service vs churn
ci = pd.crosstab(index=[df['Contract'], df['InternetService']], columns=df['Churn'])
ci.plot(kind='bar', stacked=True, figsize=(8,6))
plt.title("Contract Type and Internet Service vs. Churn", fontweight='bold')
plt.xlabel("Contract Type + Internet Service")
plt.ylabel("Proportion")
plt.legend(title='Churn', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```



- Customers with **month-to-month contracts and fiber optic internet** have the highest churn rate.
- On the other hand, customers with **two-year contracts** and either **DSL** or **no internet service** churn the least.
- This makes sense, as longer contracts and simpler services lead to customer loyalty, while high-cost, short-term plans result in more churn.

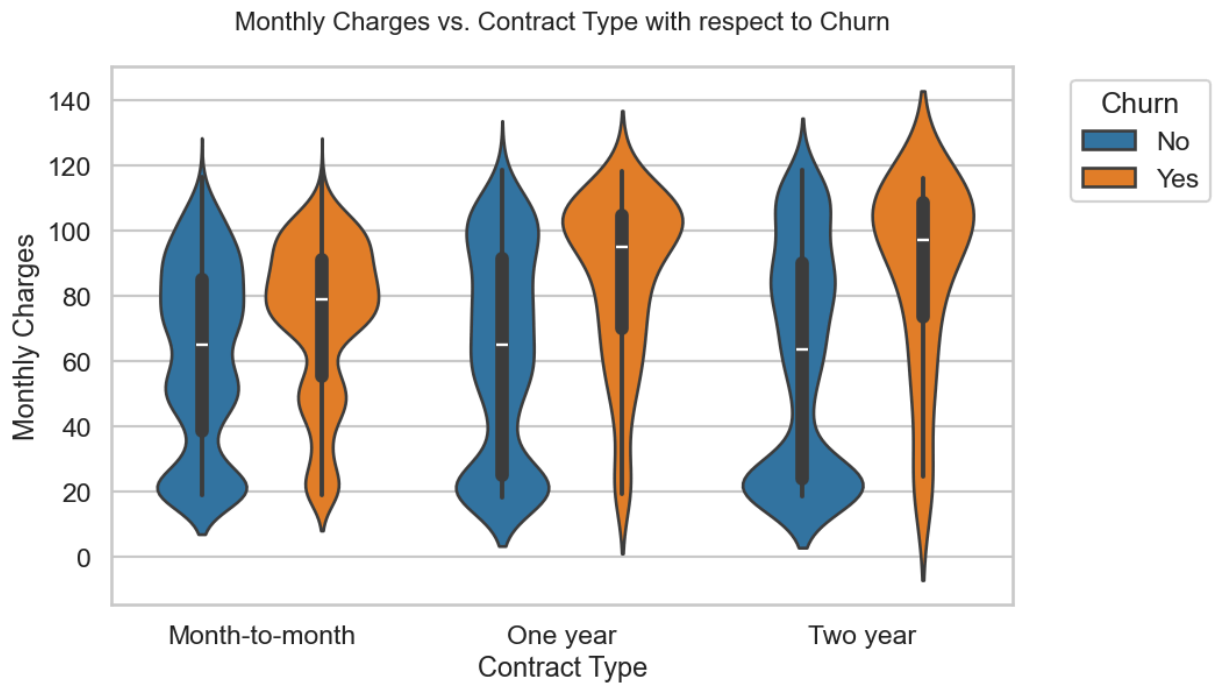
```
In [62]: # tenure + payment method vs churn
plt.figure(figsize=(10,6))
sns.boxplot(df, x='PaymentMethod', y='tenure', hue='Churn', showfliers=False)
plt.title('Tenure vs. Payment Method with respect to Churn', fontweight='bold')
plt.xlabel('Payment Method')
plt.ylabel('Tenure (Months)')
plt.xticks(rotation=45)
plt.legend(title='Churn', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```



- Customers using **electronic checks** have higher churn rates and generally lower tenure.
- People paying with **mailed checks** also leave earlier, but it's not as extreme.
- **Automatic payments** (bank transfer or credit card) are clearly associated with longer customer retention.
- Longer tenure usually means more loyalty, and automatic payment methods seem to encourage that.

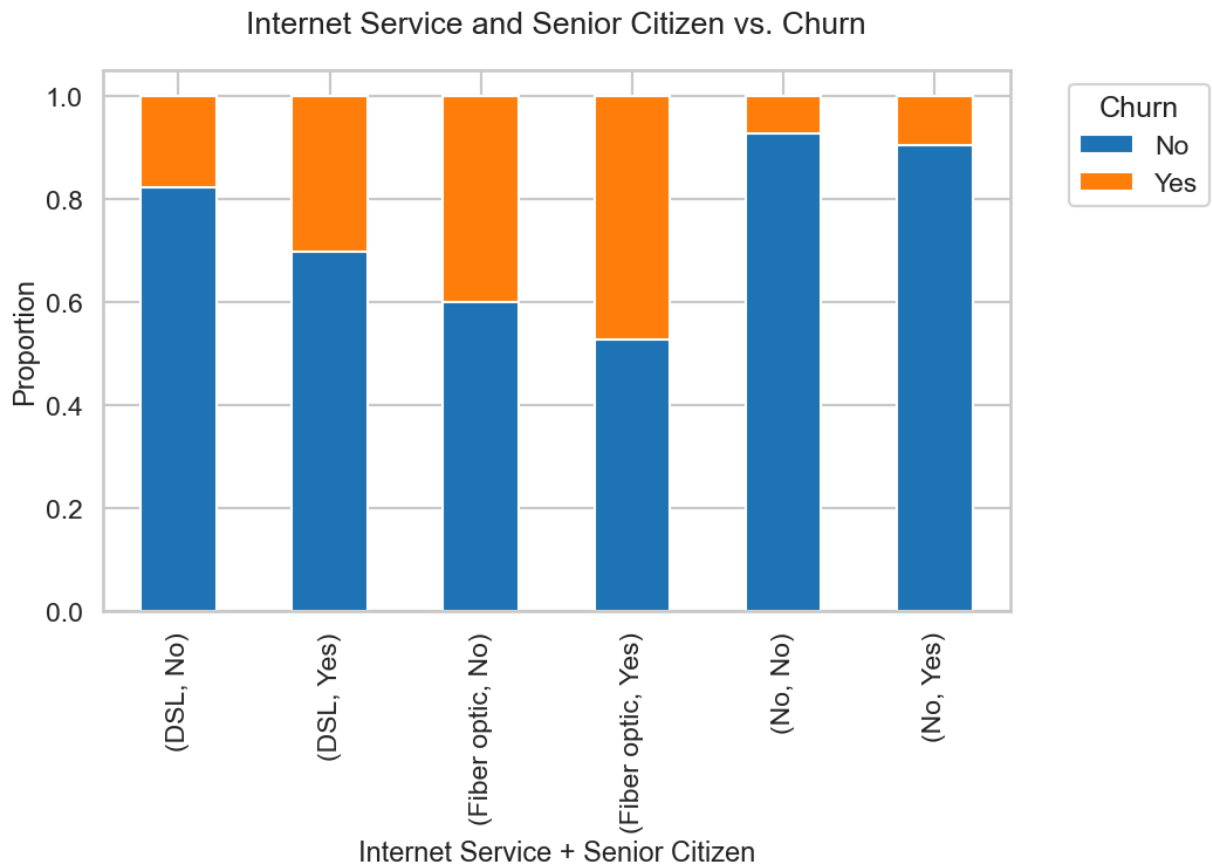
```
In [64]: # monthly charges + contract vs. churn
plt.figure(figsize=(10,6))
sns.violinplot(df, x='Contract', y='MonthlyCharges', hue="Churn")
plt.title('Monthly Charges vs. Contract Type with respect to Churn', fontsize=12)
plt.xlabel("Contract Type")
```

```
plt.ylabel("Monthly Charges")
plt.legend(title= "Churn", bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```



- Churn is highest among **month-to-month** contract holders, especially those with higher monthly charges.
- **One-year** and **two-year** contract customers have much lower churn, and their charges seem more balanced.
- Customers locked into long-term contracts are clearly more stable and less likely to leave.
- Overall, flexibility in contracts seems to lead to higher churn.

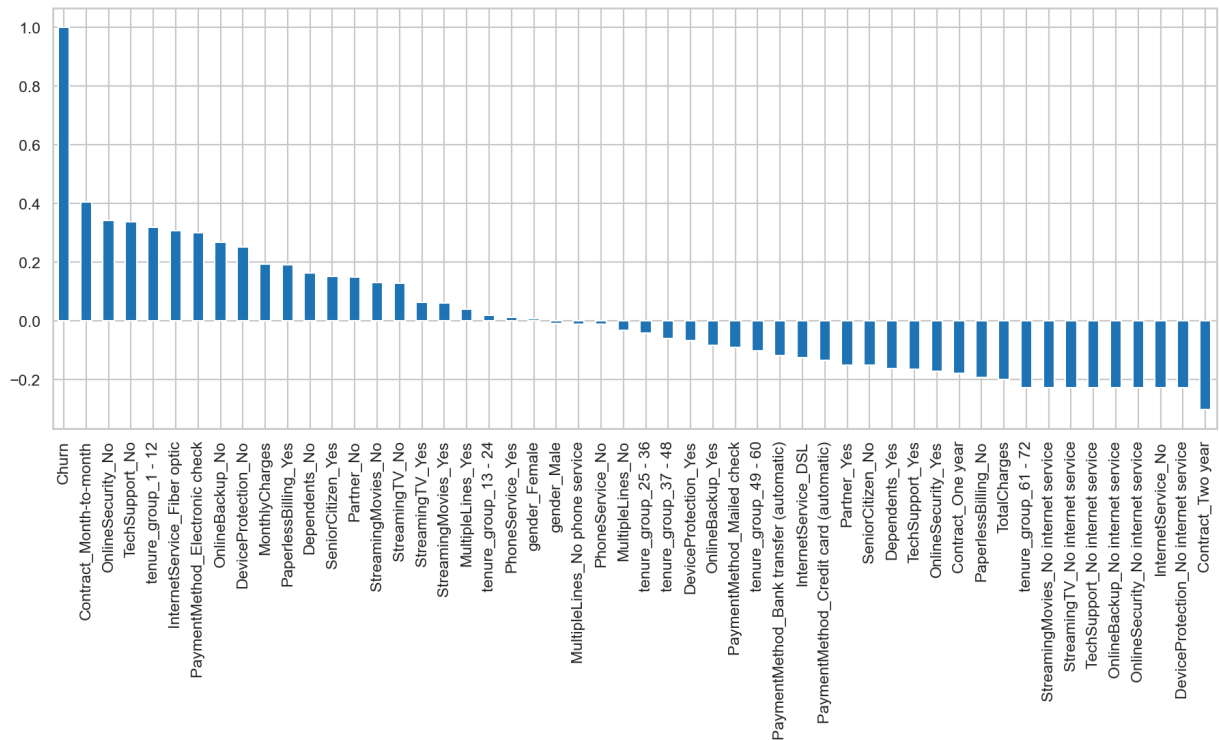
```
In [66]: # Internet Service + Senior Citizen vs Churn
its = pd.crosstab(index=[df['InternetService'], df["SeniorCitizen"]], columns=
its.plot(kind='bar', stacked=True, figsize=(10,6))
plt.title("Internet Service and Senior Citizen vs. Churn")
plt.xlabel('Internet Service + Senior Citizen')
plt.ylabel("Proportion")
plt.legend(title="Churn", bbox_to_anchor=(1.05,1), loc='upper left')
plt.show()
```



- Senior citizens with **fiber optic internet** show the highest churn rate.
- Even non-senior fiber optic users have a higher churn rate compared to those using DSL.
- **DSL users** appear to be more loyal and stable.
- Customers without internet service almost never churn, probably because they use fewer services and have fewer frustrations.
- Fiber optic + older customers = biggest churn risk zone.

```
In [68]: temp = df.copy()
temp['Churn'] = np.where(temp.Churn == 'Yes', 1, 0)
temp = temp.drop('tenure', axis=1)
```

```
In [69]: plt.figure(figsize=(22, 8))
temp1 = pd.get_dummies(temp, dtype=np.int32)
temp1.corr()['Churn'].sort_values(ascending=False).plot(kind='bar')
plt.show()
```



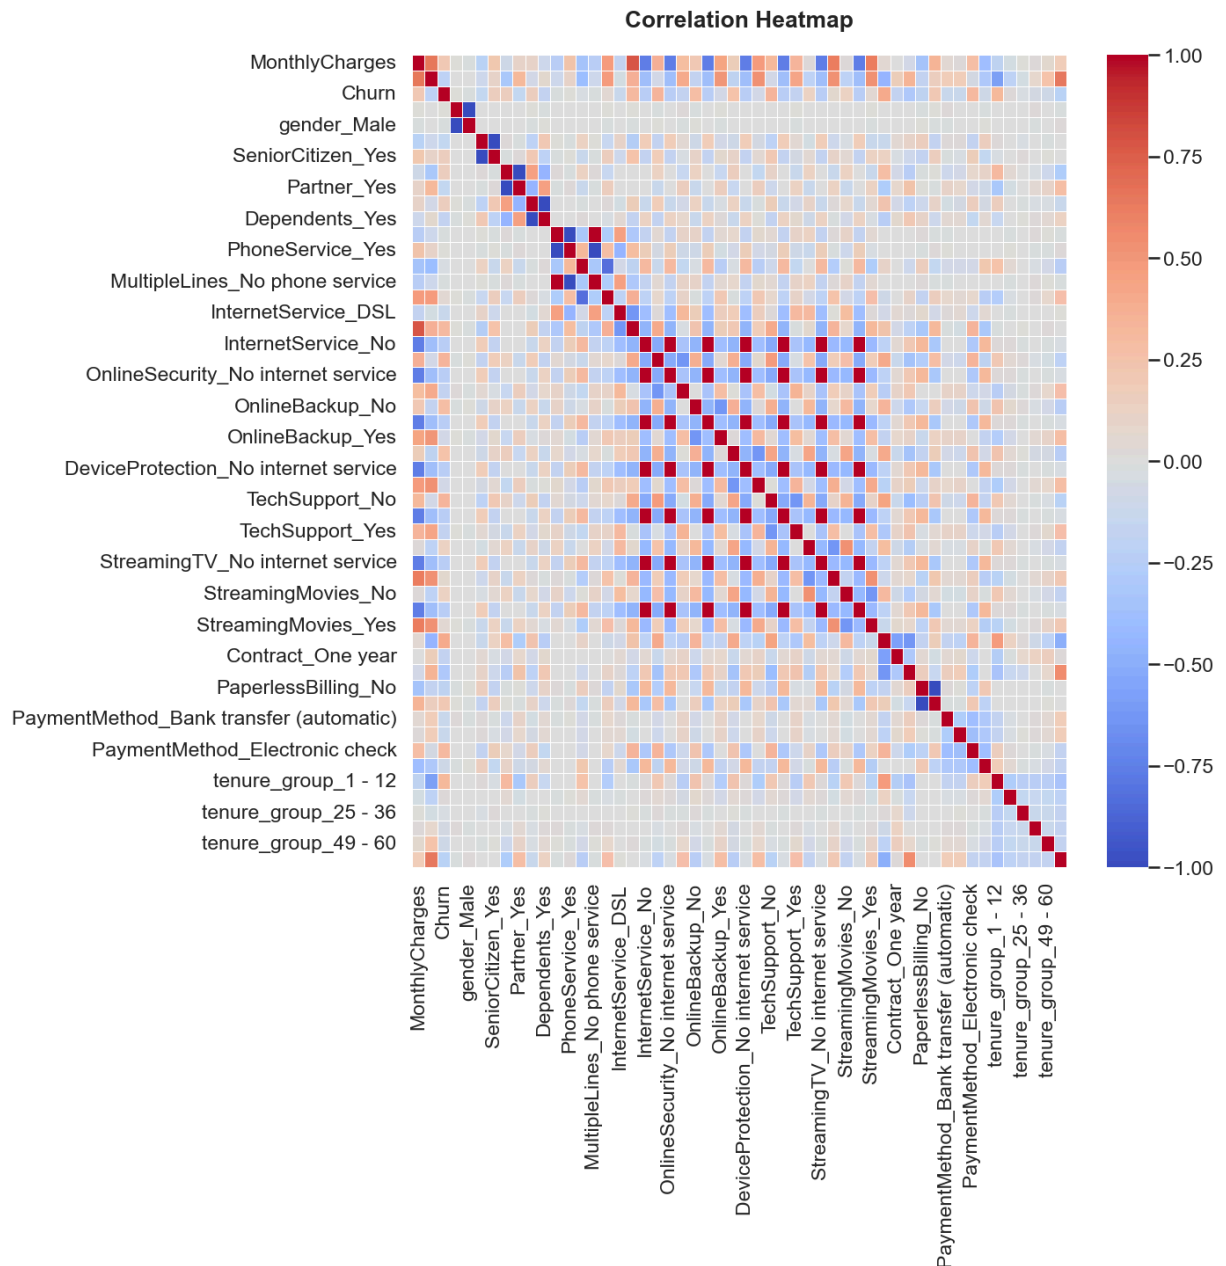
## Key Observations:

- **More churn** seen in:
  - **Fiber optic users**
  - **Electronic check payments**
  - **High monthly charges** and **paperless billing**
- **Less churn** seen in:
  - **Long-term customers (5+ years)**
  - **Two-year contracts**
  - Customers with **tech support, security, and backup services**
  - **Auto credit card payments**
- **No real impact** from:
  - **Gender, phone service, and multiple lines**

## In short:

People with high bills and fiber internet leave more. Long contracts, loyalty, and extra services help keep customers.

```
In [71]: plt.figure(figsize=(12,12))
sns.heatmap(temp1.corr(), cmap='coolwarm', linewidths=0.5)
plt.title("Correlation Heatmap", fontweight='bold')
plt.show()
```



## 🎯 What I Found from the Telco Churn Data

After going through the Telco churn data, here's what I noticed:

### 1. Churn Rate:

- About **26.5%** customers have left — that's pretty high and something to worry about.

### 2. Who's Leaving More:

- **Senior citizens** are more likely to churn.
- **Gender** doesn't seem to matter much.

### 3. Tenure Insight:



- People with **less than a year** with the company leave the most.
- Long-time customers usually stay loyal.

#### 4. **Contract Type:**

- **Month-to-month** plans have the highest churn.
- **One or two-year contracts** are much more stable.

#### 5. **Internet Service:**

- **Fiber optic** customers leave more compared to DSL or no-internet users. Probably due to cost or service issues.

#### 6. **Charges Factor:**

- Customers paying **higher monthly charges** churn more.
- But those with **high total charges** (been around longer) usually stay.

#### 7. **Payments:**

- **Electronic checks** are linked to more churn.
- Auto-pay methods seem to help keep customers.

#### 8. **Add-ons:**

- People without **security, backup, or tech support** are leaving more.
- Adding these services might help reduce churn.